

1 Basic Statistical Concepts

Consider a poll with two answers, A and B, regarding political parties. Let:

- N : total number of voters,
- M : number of voters supporting A,
- n : size of the poll,
- X_1, X_2, \dots, X_n : responses,
- Each $X_i \in \{0, 1\}$ if $X_i = 1$ supports A.

Additionally, assume:

- We select n individuals from N at random and record their truthful reply,
- Every person asked replies (no selection bias),
- People can be asked repeatedly.

The aim of the poll is to estimate the fraction of party A supporters, say θ .

Definition 1 (Estimator). *An intuitive estimator is:*

$$\hat{\theta} = \frac{1}{n} \sum_{i=1}^n X_i$$

This estimator will be analyzed in the following sections to determine whether it is unbiased, consistent, and optimal.

2 Statistical Models

Let (X, \mathcal{F}) be a measurable space, i.e., a set X with a sigma-algebra \mathcal{F} , in which our statistical observations take values.

Definition 2 (Statistical Model). *Let (X, \mathcal{F}) be some sample space. We call the parameter space Θ . A statistical model is a family of probability measures $\{P_\theta\}_{\theta \in \Theta}$.*

Remark 1. *Often (X, \mathcal{F}) is a product space. For example, if $X_i \in \{0, 1\}$, each P_θ is a product distribution, i.e., X_1, X_2, \dots, X_n are independent and identically distributed (iid). Then we say $\{P_\theta : \theta \in \Theta\}$ is an iid statistical model.*

Remark 2. *If every person could only be asked once, we would have P_θ as a hypergeometric distribution, which converges to the Bernoulli model as $N, M \rightarrow \infty$.*

3 Parameter Estimation

Assume $(\Omega, \mathcal{F}, P_\theta)$ is the setting of parametric statistics. Assume Θ is measurable.

Definition 3 (Estimator). *An estimator for θ is any measurable function $\hat{\theta} : X \rightarrow \Theta$, i.e., any function that, based on some data X , outputs a guess $\hat{\theta}(X)$ for θ .*

4 Unbiased and Consistent Estimators

4.1 Unbiased Estimator

Definition 4 (Unbiased Estimator). *Let $(\Omega, \mathcal{F}, P_\theta)$ be a measurable space. An estimator $\hat{\theta}$ is called unbiased if:*

$$\mathbb{E}[\hat{\theta}] = \theta \quad \forall \theta \in \Theta$$

where \mathbb{E}_{P_θ} denotes expectation under the law P_θ . In more explicit terms, unbiasedness means no systematic error.

Proof. For the Bernoulli model, we compute:

$$\mathbb{E}[\hat{\theta}_n] = \mathbb{E}\left[\frac{1}{n} \sum_{i=1}^n X_i\right] = \frac{1}{n} \sum_{i=1}^n \mathbb{E}[X_i] = \frac{1}{n} \sum_{i=1}^n \theta = \theta$$

Thus, $\hat{\theta}_n$ is an unbiased estimator of θ . □

4.2 Consistent Estimator

Definition 5 (Consistent Estimator). *Let $\{P_{\theta,n} : n \geq 1\}$ be a sequence of statistical models on the same parameter space. Let $\hat{\theta}_n$ be a sequence of estimators. The sequence $\hat{\theta}_n$ is called consistent if for every $\theta \in \Theta$:*

$$\hat{\theta}_n \rightarrow \theta \quad \text{in probability as } n \rightarrow \infty$$

or equivalently:

$$P_\theta \left(\lim_{n \rightarrow \infty} \hat{\theta}_n = \theta \right) = 1$$

Proof. For the Bernoulli model:

$$\hat{\theta}_n = \frac{1}{n} \sum_{i=1}^n X_i$$

We know $\mathbb{E}[\hat{\theta}_n] = \theta$ and $\text{Var}(\hat{\theta}_n) = \frac{\theta(1-\theta)}{n}$. Using Chebyshev's inequality, for any $\epsilon > 0$:

$$P \left(|\hat{\theta}_n - \theta| > \epsilon \right) \leq \frac{\text{Var}(\hat{\theta}_n)}{\epsilon^2} = \frac{\theta(1-\theta)}{n\epsilon^2}$$

As $n \rightarrow \infty$, this probability tends to 0, proving that $\hat{\theta}_n$ is consistent. □

5 Maximum Likelihood Estimation (MLE)

Definition 6 (Maximum Likelihood Estimator). *The maximum likelihood estimator (MLE) is the parameter that maximizes the likelihood function:*

$$L(\theta) = \prod_{i=1}^n P_\theta(X_i)$$

5.1 Proof: MLE for Bernoulli Model

Proof. For the Bernoulli model, $P_\theta(X_i) = \theta^{X_i}(1-\theta)^{1-X_i}$, so the likelihood function is:

$$L(\theta) = \prod_{i=1}^n \theta^{X_i}(1-\theta)^{1-X_i} = \theta^{\sum X_i} (1-\theta)^{n-\sum X_i}$$

Taking the logarithm:

$$\log L(\theta) = \sum X_i \log \theta + (n - \sum X_i) \log(1 - \theta)$$

Setting the derivative with respect to θ equal to 0 gives:

$$\frac{d}{d\theta} \log L(\theta) = \frac{\sum X_i}{\theta} - \frac{n - \sum X_i}{1 - \theta} = 0$$

Solving for θ , we get:

$$\hat{\theta}_n = \frac{1}{n} \sum_{i=1}^n X_i$$

which is the MLE. □

6 Bayesian Methods

Definition 7 (Posterior Distribution in Bayesian Inference). *In Bayesian statistics, a key element is the prior distribution, denoted by $\pi(\theta)$, which reflects our beliefs about the parameter θ before observing data. The posterior distribution is given by:*

$$\pi(\theta|X) \propto P_\theta(X)\pi(\theta)$$

6.1 Example: Posterior for Bernoulli Model

Example 1. Suppose we have a Beta prior for θ , $\pi(\theta) \sim \text{Beta}(\alpha, \beta)$, and observe X_1, \dots, X_n as Bernoulli trials. The likelihood is:

$$P(X|\theta) = \theta^{\sum X_i} (1 - \theta)^{n - \sum X_i}$$

The posterior is proportional to the product of the prior and likelihood:

$$\pi(\theta|X) \propto \theta^{\sum X_i + \alpha - 1} (1 - \theta)^{n - \sum X_i + \beta - 1}$$

Thus, $\pi(\theta|X) \sim \text{Beta}(\sum X_i + \alpha, n - \sum X_i + \beta)$.

Notes on Bayes and Posterior

Posterior = prior \times likelihood

Normalizing Constant

$$\int \text{Posterior } dx = 1$$

So,

$$\int \text{Posterior } dx = 1$$

Prior \rightarrow Posterior via Bayes.

Let \mathcal{F}_0 be a σ -algebra on Ω and suppose $(\Omega, \mathcal{F}_0, P_\theta)$ is a dominated statistical model with densities $p(x|\theta)$. Assume

$$x, \theta \in \Omega \Rightarrow p(x|\theta)$$

is jointly measurable with respect to $\mathcal{F}_0 \times \mathcal{F}_1$.

Let π be a prior distribution on Ω with density $\pi(\theta)$ with respect to measure ν . Define posterior density

$$\pi(\theta|x) = \frac{p(x|\theta)\pi(\theta)}{\int p(x|\theta)\pi(\theta) d\theta}$$

The corresponding probability measure is called the **posterior distribution**.

Think of $p(x|\theta)$ as a Lebesgue measure. Let ν be a Lebesgue density.

Exception: If $\Omega = \{0, 1\}$, then we take ν to be the counting measure.

From the posterior, we can derive several estimators. For example, $E[\theta|X = x]$ is convex:

$$\int \theta p(x|\theta) d\theta = E[\theta|X = x]$$

Example: Binomial model $X|\theta \sim \text{Binomial}(n, \theta)$ with prior $\theta \sim \text{Unif}(0, 1)$.

For a uniform prior, we know the MAP and MLE.

Posterior mean:

$$\theta_{\text{MAP}} = \frac{k+1}{n+2}$$

In the case of coin flips, $X \sim \text{Binomial}(n, \theta)$, where k is the number of heads, we conclude $\theta|X \sim \text{Beta}(k+1, n-k+1)$.

$$\theta|X \sim \text{Beta}(k+1, n-k+1)$$

Conjugate Bayes Models: Let $P_\theta \in \mathcal{P}$ be a statistical model. Then some family of priors is called **conjugate** if

$$P_\theta \in \mathcal{P} \Rightarrow \theta|X \in \mathcal{P}$$

for all $X \in \mathcal{X}$, where \mathcal{X} is the sample space.

$$\theta|X \sim \text{Beta}(a, b), \quad X \sim \text{Bernoulli}(p)$$

Loss Functions and Risk

Loss Function: A function $L : \Theta \times \mathcal{X} \rightarrow [0, \infty)$ is a basis function if for every $\theta \in \Theta$, $L(\theta, \cdot)$ is measurable.

Given an estimator δ , the expected loss is

$$R(\theta, \delta) = E_\theta[L(\theta, \delta)]$$

Mean Squared Error (MSE):

$$L(x, y) = (x - y)^2 \Rightarrow R(\theta, \delta) = E_\theta[(\delta - \theta)^2]$$

Bias-Variance Decomposition:

$$L(x, y) = (x - y)^2$$

Proof: Let $\delta(x) = E[\theta|X = x]$.

$$R(\theta, \delta) = E_\theta[(\delta(X) - \theta)^2]$$

Bias-variance decomposition:

$$E[(\delta(X) - \theta)^2] = \text{Var}(\delta(X)) + (\text{Bias})^2$$

Minimax and Bayes Risk

Minimax Risk: Given an estimator δ in a model $P_\theta \in \mathcal{P}$, the maximal risk of it is

$$\sup_{\theta \in \Theta} R(\theta, \delta)$$

The minimax of a model P_θ is given as $\inf_\delta \sup_\theta R(\theta, \delta)$, where the inf is over all estimators.

An estimator is called minimax if

$$\sup_\theta R(\theta, \delta) = \inf_\delta \sup_\theta R(\theta, \delta)$$

Bayes Risk: Given an estimator δ and prior π on Θ , the Bayes risk of δ is defined as

$$R_\pi(\delta) = \int R(\theta, \delta) d\pi(\theta)$$

The posterior risk of an estimator $\delta(X)$ is defined by

$$R(\delta|X = x) = E[L(\theta, \delta(X))|X = x]$$

Suppose δ^* is an estimator that minimizes the posterior risk, $\delta^*(x) = E[\theta|X = x]$. Then it also minimizes the Bayes risk. If $L(x, y) = (x - y)^2$, the Bayes optimal estimator $\delta(x)$ is the posterior mean.

We want to construct $C(x)$ s.t. $P_\theta(\theta \in C(x)) \geq 1 - \alpha, \forall \theta \in [0, 1]$

$$x^{(1)} \quad (\quad) \quad C(x^{(1)})$$

$$x^{(k)} \quad (\quad) \quad C(x^{(k)})$$

$$\theta \rightarrow \quad \rightarrow \quad \rightarrow \quad \text{contains true param } 3/4 \text{ times}$$

Example cont.:

Best guess: $C(x) = \left[\frac{\bar{X}_n - a}{n}, \frac{\bar{X}_n + b}{n} \right]$

$$P_\theta^n(\theta \in C(x)) = P_\theta^n \left(\frac{\bar{X}_n}{n} - \theta \in [-b, a] \right)$$

$$= F_\theta^n(a) - F_\theta^n(-b) + \rho_n$$

where $F_\theta^n : \mathbb{R} \rightarrow [0, 1]$, $F_\theta^n(t) = P_\theta^n \left(\frac{\bar{X}_n - \theta}{n} \leq t \right)$ is the CDF of $\frac{\bar{X}_n - \theta}{n}$ under P_θ and $\rho_n = P_\theta^n \left(\frac{\bar{X}_n}{n} - \theta = -b \right)$.

How to choose a and b:

$$\text{CDF} \quad \text{CDF} \quad \leftarrow \quad -b \quad a \rightarrow t$$

We'd like to choose $a = (F_\theta^n)^{-1} \left(1 - \frac{\alpha}{2} \right)$ and $b = (F_\theta^n)^{-1} \left(\frac{\alpha}{2} \right)$, where

$$(F_\theta^n)^{-1}(p) := \inf \{ t \in \mathbb{R} : F_\theta^n(t) \geq p \} \quad (\text{Quantile Function})$$

Let's use a normal approximation, for $\sigma^2 = \theta(1 - \theta)$:

$$\sqrt{n} \left(\frac{\bar{X}_n}{n} - \theta \right) = \frac{1}{\sqrt{n}} \sum_{k=1}^n \frac{X_k - \theta}{\sigma} \xrightarrow{d} \mathcal{N}(0, 1) \quad [\text{CLT}]$$

$$X_k \sim \text{Ber}(\theta)$$

Then it follows that

$$F_\theta^n(a_n) = P_\theta^n \left(\frac{\bar{X}_n}{n} - \theta \leq a_n \right)$$

$$= P_\theta^n \left(\frac{\sqrt{n}}{\sigma} \left(\frac{\bar{X}_n - \theta}{n} \right) \leq \sqrt{n} a_n \right)$$

$$= \Phi \left(\frac{\sqrt{n}}{\sigma} a_n \right),$$

where the convergence is valid if $a_n := \text{const.} \cdot \frac{1}{\sqrt{n}}$.

Now, let us choose

$$a := \frac{\sigma}{\sqrt{n}} z_{1 - \frac{\alpha}{2}}$$

where $z_{1 - \frac{\alpha}{2}} = \Phi^{-1} \left(1 - \frac{\alpha}{2} \right)$ is the $1 - \frac{\alpha}{2}$ quantile of $\mathcal{N}(0, 1)$ and $b = a$. Then

$$C(x) = \left[\frac{\bar{X}_n}{n} - \frac{\sigma}{\sqrt{n}} z_{1 - \frac{\alpha}{2}}, \frac{\bar{X}_n}{n} + \frac{\sigma}{\sqrt{n}} z_{1 - \frac{\alpha}{2}} \right]$$

It follows

$$P_\theta^n(\theta \in C(x)) = F_\theta^n(a_n) - F_\theta^n(b) + \rho_n = 1 - \frac{\alpha}{2} + o(1) + o(1)$$

$$= 1 - \alpha + o(1) \text{ as } n \rightarrow \infty$$

\Rightarrow Asymptotically valid confidence set

One more problem: σ depends on θ

- Upper bound: $\sup_{\theta \in [0, 1]} \theta(1 - \theta) = \frac{1}{4}$ (maximized at $\theta = \frac{1}{2}$)
- Empirical Variance: $\hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (X_i - \frac{1}{n} \sum_{i=1}^n X_i)^2$

$$\frac{\hat{\sigma}^2}{\sigma^2} \xrightarrow{P_\theta} 1$$

Slutsky's Theorem:

$$X_n \xrightarrow{d} X, \quad Y_n \xrightarrow{d} \text{const.} \Rightarrow X_n Y_n \xrightarrow{d} CX$$

Exercise: Use this to deduce that $a_n = \frac{\hat{\sigma}}{\sqrt{n}} z_{1-\frac{\alpha}{2}}$ is also valid

Remark:**Hypothesis Testing**

Definition: Let $(P_\theta : \theta \in \Theta)$ be a statistical model and let $\Theta = \Theta_0 \cup \Theta_1$ be a partition. Then:

- A statistical test is a measurable function of the data $\varphi : (\mathcal{X}, \mathcal{F}) \rightarrow [0, 1]$
- If $\forall x \in \mathcal{X}, \varphi(x) \in \{0, 1\}$, then φ is a non-randomized test
- Else φ is randomized

Definitions:

- $H_0 : \theta \in \Theta_0$ is called the null hypothesis
- $H_1 : \theta \in \Theta_1$ is called the alternative hypothesis
- The map $\theta \rightarrow \beta_\varphi(\theta) = P_\theta[\varphi = 1]$ is called the power function of a test φ

$$\begin{array}{ccccc} 1 & \beta_\varphi(\theta) & 0 & \Theta_0 & \Theta_1 & \Theta \end{array}$$

- For $\theta \in \Theta_0$, $\beta_\varphi(\theta)$ is the type-I-error under θ [Wrongly rejecting the null]
- For $\theta \in \Theta_1$, $1 - \beta_\varphi(\theta)$ is the type-II-error

Note:

$$1 - P_\theta(\varphi = 1) = P_\theta(\varphi = 0) = P_\theta(\text{wrongly accepting the null})$$

Definition: [Level]

$\varphi : \mathcal{X} \rightarrow [0, 1]$ has level $\alpha \in [0, 1]$ if

$$\sup_{\theta \in \Theta_0} \beta_\varphi(\theta) \leq \alpha$$

Definition: [Uniformly most powerful test]

Given a level $\alpha \in (0, 1)$, $\varphi : \mathcal{X} \rightarrow [0, 1]$ is called UMP if for every other test φ' of level α and all $\theta \in \Theta_1$,

$$\beta_\varphi(\theta) \geq \beta_{\varphi'}(\theta)$$

$$\begin{array}{ccccc} 1 & \alpha & 0 & \beta_\varphi(\theta) & \beta_{\varphi'}(\theta) & \Theta_0 & \Theta_1 \end{array}$$

Remark:

In general, it is very hard to find UMP tests. But: for simple hypotheses, i.e. $\Theta_0 = \{\theta_0\}, \Theta_1 = \{\theta_1\}$, it is possible. Here, likelihood ratio tests are UMP.

Theorem: [Neyman-Pearson Lemma]

Let $\Theta_0 = \{\theta_0\}, \Theta_1 = \{\theta_1\}$ be simple:

1. **Existence:** There exists a test φ and a constant $k \in [0, \infty)$, s.t. $P_{\theta_0}(\varphi = 1) = \alpha$, of the form

$$\varphi(x) = \begin{cases} 1, & \text{if } \frac{p_{\theta_1}(x)}{p_{\theta_0}(x)} > k \\ 0, & \text{if } \frac{p_{\theta_1}(x)}{p_{\theta_0}(x)} < k \end{cases} \quad (*)$$

Here $p_{\theta_1}, p_{\theta_0}$ are densities w.r.t. some dominated measure μ , e.g. $\mu = p_{\theta_0} + p_{\theta_1}$. Finite Θ implies measure is always dominated (likelihood always exists).

2. **Sufficiency:** If φ satisfies $P_{\theta_0}(\varphi = 1) = \alpha$ and $(*)$ then φ is a UMP level α test.
3. **Necessity:** If φ_k is UMP for level α , then it must be of the form $(*)$, and it also satisfies $P_{\theta_0}(\varphi_k = 1) = \alpha$, or else it must satisfy $P_{\theta_1}(\varphi_k = 1) = 1$.

Proof:

1. Define $r(x) = \frac{p_{\theta_1}(x)}{p_{\theta_0}(x)} \in [0, \infty) \cup \{\pm\infty\}$. Let F_0 be the CDF of $r(x)$ under P_{θ_0} .

$$F_0(t) = P_{\theta_0}(r(x) \leq t)$$

Then define also $\alpha(t) = 1 - F_0(t) = P_{\theta_0}(r(x) > t)$

- α is right-continuous:

$$\lim_{\epsilon \rightarrow 0} \alpha(t + \epsilon) = \lim_{\epsilon \rightarrow 0} P_{\theta_0}(r(x) > t + \epsilon) = P_{\theta_0}(r(x) > t) = \alpha(t)$$

- α is non-increasing
- α has left limits

$$\lim_{\epsilon \rightarrow 0} \alpha(t - \epsilon) = P_{\theta_0}(r(x) > t - \epsilon) = \alpha(t^-)$$

α is **cadlag**:

- Continuous from the right
- Limit from the left

There exists some $k \in [0, \infty)$ s.t. $\alpha \leq \alpha(k^-)$ and $\alpha \geq \alpha(k)$

We define our test

$$\varphi(x) = \begin{cases} 1 & \text{if } r(x) > k \\ \gamma & \text{if } r(x) = k \quad [\text{reject null w.p. } \gamma] \\ 0 & \text{if } r(x) < k \end{cases}$$

We set

$$\gamma = \frac{\alpha - \alpha(k)}{\alpha(k^-) - \alpha(k)}$$

The level of φ is

$$\begin{aligned} E_{\theta_0}[\varphi(x)] &= P_{\theta_0}(\varphi(x) = 1) \\ &= P_{\theta_0}(r(x) > k) + P_{\theta_0}(r(x) = k) \cdot \gamma \\ &= \alpha(k) + [\alpha(k^-) - \alpha(k)] \cdot \frac{\alpha - \alpha(k)}{\alpha(k^-) - \alpha(k)} = \alpha \\ &\quad (\text{randomizing the test}) \end{aligned}$$

Lecture 6

Neyman-Pearson

Power of a test:

$$E_{\theta_1}[\varphi] = P_{\theta_1}(\varphi = 1)$$

Likelihood ratio test:

$$\frac{p_{\theta_1}(x)}{p_{\theta_0}(x)} = r(x)$$

LR test

$$\varphi(x) = \begin{cases} 1 & \text{if } r(x) > k \\ \gamma & \text{if } r(x) = k \\ 0 & \text{if } r(x) < k \end{cases}$$

for some $k \in [0, \infty)$, $\gamma \in [0, 1]$.

Note: LR tests are UMP for simple hypothesis testing:

- Given some α , if LR satisfies $E_{\theta_0}[\varphi] = \alpha$, it represents a Type I error.
- φ minimizes the Type II error

$$E_{\theta_1}[\varphi] \geq E_{\theta_1}[\varphi'] \quad \forall \varphi'$$

Cont. of proof (part of UMP)

Let φ' be another level α test, $E_{\theta_0}[\varphi'] \leq \alpha$.

Goal: $E_{\theta_1}[\varphi] \geq E_{\theta_1}[\varphi']$. Let μ be the dominating measure.

Consider

$$\int (\varphi(x) - \varphi'(x))(p_{\theta_1}(x) - kp_{\theta_0}(x)) d\mu(x) = 0$$

Claim: $p \geq 0$.

Observe:

- If $p_{\theta_1}(x) - kp_{\theta_0}(x) > 0 \Rightarrow \frac{p_{\theta_1}(x)}{p_{\theta_0}(x)} > k \Rightarrow \varphi(x) = 1$.
- If $p_{\theta_1}(x) - kp_{\theta_0}(x) < 0 \Rightarrow \varphi(x) = 0$.
- If $p_{\theta_1}(x) - kp_{\theta_0}(x) = 0 \Rightarrow \text{integrand} = 0$.

$$\Rightarrow p = 0$$

$$\Rightarrow \int (\varphi - \varphi') p_{\theta_1} d\mu = \int (\varphi - \varphi') p_{\theta_0} d\mu = k [E_{\theta_0}[\varphi] - E_{\theta_0}[\varphi']] \geq 0$$

$$\Rightarrow E_{\theta_1}[\varphi] \geq E_{\theta_1}[\varphi']$$

Part (3) UMP \Rightarrow (LR): Take φ^* a UMP test, $E_{\theta_0}[\varphi^*] = \alpha$, and let φ be the LR test with $E_{\theta_0}[\varphi] = \alpha$ with (*).

Goal: $\varphi = \varphi^*$ a.e. except on $\{r(x) = k\}$.

Define

$$x^+ = \{x : \varphi(x) > \varphi^*(x)\}$$

$$x^- = \{x : \varphi(x) < \varphi^*(x)\}$$

$$x^0 = \{x : \varphi(x) = \varphi^*(x)\}$$

$$\tilde{x} = (x^+ \cup x^-) \cap \{x : p_{\theta_1}(x) \neq kp_{\theta_0}(x)\}$$

It suffices to show $\mu(\tilde{x}) = 0$.

Like before, we have

$$(\varphi - \varphi^*)(p_{\theta_1} - kp_{\theta_0}) > 0 \text{ on } \tilde{x}$$

Thus if $\mu(\tilde{x}) > 0$,

$$\begin{aligned} \int_{\mathcal{X}} (\varphi - \varphi^*)(p_{\theta_1} - kp_{\theta_0}) d\mu &\geq 0 \\ \int_{\tilde{x}} (\varphi - \varphi^*)(p_{\theta_1} - kp_{\theta_0}) d\mu &\geq 0 \end{aligned}$$

But also

$$E_{\theta_1}[\varphi] - E_{\theta_1}[\varphi^*] > k [E_{\theta_0}[\varphi] - E_{\theta_0}[\varphi^*]] \geq 0$$

\Rightarrow Cannot be φ^* is UMP.

Example (Gaussian Location Model)

$$X_1, \dots, X_n \stackrel{\text{iid}}{\sim} \mathcal{N}(\mu, \sigma^2)$$

$$H_0 : \mu = \mu_0, \quad H_1 : \mu = \mu_1, \quad \mu_0 < \mu_1$$

Then:

$$\begin{aligned} \frac{p_1(X_1, \dots, X_n)}{p_0(X_1, \dots, X_n)} &= \exp \left(-\frac{1}{2\sigma^2} \sum_{i=1}^n (X_i - \mu_1)^2 + \frac{1}{2\sigma^2} \sum_{i=1}^n (X_i - \mu_0)^2 \right) \\ &= \exp \left(-\frac{1}{2\sigma^2} \sum_{i=1}^n (\mu_1^2 - \mu_0^2) - \frac{2(\mu_1 - \mu_0)}{\sigma^2} \sum_{i=1}^n X_i \right) \\ &= \exp \left(-\frac{n}{2\sigma^2} (\mu_1^2 - \mu_0^2) - \frac{2(\mu_1 - \mu_0)}{\sigma^2} \sum_{i=1}^n X_i \right) \geq K_\alpha \\ &\Rightarrow \frac{1}{n} \sum_{i=1}^n X_i \geq K_\alpha, \text{ some } K_\alpha \in \mathbb{R} \end{aligned}$$

To determine K_α :

$$\begin{aligned} \bar{X}_n &:= \frac{1}{n} \sum_{i=1}^n X_i \stackrel{H_0}{\sim} \mathcal{N}(\mu_0, \sigma^2/n) \\ \Rightarrow \mathbb{L} &= P_{H_0}(\bar{X}_n \geq K_\alpha) = 1 - P_{H_0}(\bar{X}_n < K_\alpha) \\ &= 1 - \Phi \left(\frac{\sqrt{n}}{\sigma} (K_\alpha - \mu_0) \right) \quad (\text{CDF for } \mathcal{N}(0, 1)) \\ \Rightarrow \text{solving for } K_\alpha &\text{ gives } K_\alpha = \mu_0 + \frac{\sigma}{\sqrt{n}} \Phi^{-1}(1 - \alpha), \\ \varphi(X_1, \dots, X_n) &= \begin{cases} 1 & \text{if } \bar{X}_n \geq \mu_0 + \frac{\sigma}{\sqrt{n}} \Phi^{-1}(1 - \alpha) \\ 0 & \text{else} \end{cases} \end{aligned}$$

Corollary

Consider simple hypothesis testing. Let φ be UMP, for level α . Then,

$$\alpha = E_{H_0}[\varphi] = E_{\theta_0}[\varphi] \leq E_{\theta_1}[\varphi]$$

Suppose $E_{\theta_1}[\varphi] = E_{\theta_1}[\varphi_0]$ then φ_0 is also UMP, $\Rightarrow \varphi_0$ is an LR test.

$$\varphi_0 = \begin{cases} 1 & \text{if } \frac{p_{\theta_1}}{p_{\theta_0}} \geq K \quad \text{a.s., some } K \\ 0 & \text{if } \frac{p_{\theta_1}}{p_{\theta_0}} < K \end{cases}$$

Also since $\varphi_0 \in \{\varphi, \beta\}$ we conclude that $p_{\theta_1} = K p_{\theta_0}$ a.s.

But

$$L = \int p_{\theta_0} d\mu = K \int p_{\theta_0} d\mu = 1 \Rightarrow K = 1$$

Correspondence theorem

$$\text{Tests} \longleftrightarrow \text{Confidence regions } C(x)$$

$$\Pr_{\theta}(\theta \in C(x)) \geq 1 - \alpha$$

$$\text{If } \Pr_{\theta}(\phi_{\theta} = 1) = \alpha$$

Theorem: Let $(P_{\theta} : \theta \in \Theta)$ be a statistical model, $\alpha \in (0, 1)$.

(i) Let $C = C(X)$ be a level- α confidence set, then

$$\phi_{\theta_0}(x) = 1 \{\theta_0 \notin C(x)\}$$

is a level- α test of $\theta = \theta_0$ vs. $\theta \neq \theta_0$.

(ii) Suppose $\{\phi_{\theta_0} : \theta_0 \in \Theta\}$ is a family of level- α tests, then

$$C(X) = \{\theta \in \Theta : \phi_{\theta}(X) = 0\}$$

is a $(1 - \alpha)$ confidence set.

Proof:

$$(i) \quad \Pr_{\theta_0}(\phi_{\theta_0} = 1) = \Pr_{\theta_0}(\theta_0 \notin C(X)) = \alpha$$

$$(ii) \quad \Pr_{\theta}(\theta \notin C(X)) = \Pr_{\theta}(\theta \notin \{\tilde{\theta} \in \Theta : \phi_{\tilde{\theta}}(X) = 0\}) = \Pr_{\theta}(\phi_{\theta}(X) = 1) \leq \alpha$$

UMPT Tests in Models with Monotone Likelihoods

Proposition: Let $\Theta \subseteq \mathbb{R}$. Consider testing $H_0 : \theta \leq \theta_0$ vs. $H_1 : \theta > \theta_0$, for some $\theta_0 \in \mathbb{R}$.

Assume there exists some test statistic $T : X \rightarrow \mathbb{R}$ and a function $h : \mathbb{R} \times \Theta \times \Theta$ such that

$$\frac{P_{\theta}(X)}{P_{\tilde{\theta}}(X)} = h(T(X), \theta, \tilde{\theta})$$

and for all $\theta \geq \tilde{\theta}$, $t \mapsto h(t, \theta, \tilde{\theta})$ is monotone increasing.

The simplest model for the relationship between Y_i and X_i assumes a linear relationship:

$$Y_i = aX_i + b + \varepsilon_i$$

for $i = 1, \dots, n$, where ε_i is centered, i.e., $E(\varepsilon) = 0$ and $\text{Var}(\varepsilon) = \sigma^2$. Suppose $\varepsilon \sim N(0, \sigma^2)$ with σ known.

The statistical model is given by

$$(\mathbb{R}, B(\mathbb{R}), (\bigotimes_{i=1}^n N(ax_i + b, \sigma^2))_{(a,b) \in \mathbb{R}^2})$$

The likelihood within the statistical model is

$$L((a, b)|y) = \prod_{i=1}^n (2\pi\sigma^2)^{-1/2} \exp\left(-\frac{1}{2\sigma^2} \sum_{i=1}^n (y_i - ax_i - b)^2\right)$$

The MLE satisfies the optimization problem

$$(\hat{a}, \hat{b}) = \arg \min_{(a,b) \in \mathbb{R}^2} \sum_{i=1}^n (y_i - (ax_i + b))^2$$

Provided that $x_i \neq x_j$ for $i \neq j$, the least squares problem has a solution with minimum given by (Gauss, 1801):

$$(\hat{a}, \hat{b}) = \left(\frac{\frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2}, \bar{y} - \hat{a}\bar{x} \right)$$

Definition 8 (Linear Model). A random vector $Y = (Y_1, \dots, Y_n)^T \in \mathbb{R}^n$ stems from a linear model if there exists a parameter vector $\beta \in \mathbb{R}^p$, a matrix $X \in \mathbb{R}^{n \times p}$, and a random vector $\varepsilon \in \mathbb{R}^n$ such that

$$Y = X\beta + \varepsilon$$

1. A linear model is called regular if

- (a) $p \leq n$ (parameter size is smaller than sample size),
- (b) X has full rank. $\text{rank}(X) = p \leq n$ (design with full rank)
- (c) $E(\varepsilon) = 0$ (noise is controlled)
- (d) The covariance matrix is positive definite, $\Sigma = (\text{Cov}(\varepsilon_i, \varepsilon_j))_{i,j \in [n]}$

2. A linear model is called ordinary if $\Sigma = \sigma^2 E_n$ (and is usually the noise is Gaussian)

Remark 3. 1. There are several synonyms

- (a) Y a dependent variable, response, regressand
- (b) X , a independent variable, predictor, design matrix, regressor
- (c) ε Error, perturbation, reression function

2. The matrix Σ is symmetric and diagonalizable, i.e. $\Sigma = UDU^T$ for some diagonal matrix, $D = \text{diag}(\lambda_1, \dots, \lambda_n) \in \mathbb{R}^{n \times n}$

3. Positive semi-definite, i.e. $\lambda_i \geq 0$

$$\begin{aligned} \langle \Sigma u, u \rangle &= \langle E[(\varepsilon - E[\varepsilon])(\varepsilon - E[\varepsilon])^T] u, u \rangle \\ &= E[(\varepsilon - E[\varepsilon])^2] \geq 0, u \in \mathbb{R}^n \end{aligned}$$

item If Σ is positive definite ($\lambda_i > 0$) for $i = 1, \dots, n$, then there exists the inverse $\Sigma^{-1} = UD^{-1}U^T$ and $\Sigma^{-1/2} = UD^{-1/2}U^T$.

4. If X is not deterministic, we speak of random design.

In the regular linear model, $\hat{\beta}$ is called weighted least squares estimate, (LSE). if

$$\|\sigma^{-1/2}(Y - X\hat{\beta})\|^2 = \inf_{\beta \in \mathbb{R}^n} \|\sigma^{-1/2}(Y - X\beta)\|^2 = \inf_{\beta \in \mathbb{R}^n} \|\sigma^{-1/2}Y - X_{\Sigma}\beta\|^2$$

where $X_{\Sigma} = \Sigma^{-1/2}X$. $X_{\Sigma}\hat{\beta}$ is the point within the subspace,

$$U = \{X_{\Sigma}\beta \mid \beta \in \mathbb{R}^n\} \subseteq \mathbb{R}^n$$

with the smallest distance to the vector $\Sigma^{-1/2}Y$. Thus, $X_{\Sigma}\hat{\beta} = \Pi_U(\Sigma^{-1/2}Y)$ where Π_U is the orthogonal projection onto U . $\Pi_U u = u$ for all $u \in U$. $\langle \Pi_U v - v, u \rangle = 0$ for all $u \in U$ and $v \in \mathbb{R}^n$. Provided that $(X_{\Sigma}^T X_{\Sigma})^{-1}$ exists, we can confirm by direct computation that the projection satisfies

$$\Pi_U = X_{\Sigma}(X_{\Sigma}^T X_{\Sigma})^{-1} X_{\Sigma}^T$$

For $u = X_{\Sigma}\beta$ we have,

$$X_{\Sigma}(X_{\Sigma}^T X_{\Sigma})^{-1} X_{\Sigma}^T X_{\Sigma}\beta = X_{\Sigma}\beta = u$$

By symmetry,

$$\langle \Pi_U v - v, u \rangle = \langle v, \Pi_U u \rangle - \langle v, u \rangle = \langle v, u \rangle - \langle v, u \rangle = 0$$

for all $u \in U$.

Lemma 1. *Representation for the LSE Consider a regular linear model, then the LSE exists uniquely, and is given by*

$$\hat{\beta} = (X_{\Sigma}^T X_{\Sigma})^{-1} X_{\Sigma}^T \Sigma^{-1/2} Y = X_{\Sigma}^+ \Sigma^{-1/2} Y$$

Proof. $\ker(X_{\Sigma}^T X_{\Sigma})$ is invertible. Suppose that $X_{\Sigma}^T X_{\Sigma} v = 0$ ($v \in \ker(X_{\Sigma}^T X_{\Sigma})$)

$$0 = v^T X_{\Sigma}^T X_{\Sigma} v = (X_{\Sigma}^T v)^T X_{\Sigma} v = \langle X_{\Sigma} v, X_{\Sigma} v \rangle = \|X_{\Sigma} v\|^2 = \|\Sigma^{-1/2} X v\|^2 \implies \|X v\|^2 = 0 \implies v = 0$$

So then

$$\begin{aligned} X_{\Sigma}\hat{\beta} &= \Pi_U \Sigma^{-1/2} Y = X_{\Sigma}(X_{\Sigma}^T X_{\Sigma})^{-1} X_{\Sigma}^T \Sigma^{-1/2} Y \\ X_{\Sigma}^T X_{\Sigma}\hat{\beta} &= X_{\Sigma}^T X_{\Sigma}(X_{\Sigma}^T X_{\Sigma})^{-1} X_{\Sigma}^T \Sigma^{-1/2} Y \\ &\implies \hat{\beta} = (X_{\Sigma}^T X_{\Sigma})^{-1} X_{\Sigma}^T \Sigma^{-1/2} Y \end{aligned}$$

□

Remark 4. 1. If $p > n$, then $(X_{\Sigma}^T X_{\Sigma})^{-1}$ does not exist and the LSE is not unique.

$$\left\{ \beta \cdot \|\Sigma^{-1/2} Y - X_{\Sigma}\beta\|^2 = 0 \right\}$$

is a $p - n$ dim subspace and each solution interpolates the data

Theorem 1. *Optimality of the LSE, Gauss-Markov Theorem Consider an ordinary linear model for $\sigma > 0$, then*

1. The least squares estimator $\hat{\beta} = (X^T X)^{-1} X^T Y$ is linear and the unbiased parameter for the parameter β .
2. For the desired parameter $\alpha = \langle \beta, v \rangle$ for $v \in \mathbb{R}$, the estimator $\hat{\alpha} = \langle \hat{\beta}, v \rangle$ is the best linear unbiased estimator (BLUE), meaning that $\hat{\alpha}$ has the optimal value within the class of linear unbiased estimators for α
3. $\hat{\sigma}^2 = \frac{\|Y - X\hat{\beta}\|^2}{n-p}$ is an unbiased estimator of σ^2

Proof.

$$\hat{\beta}(y + \tilde{y}) = \hat{\beta}(y) + \hat{\beta}(\tilde{y}) \text{ for } y, \tilde{y} \in \mathbb{R}^n$$

$$E[\hat{\beta}] = (X^T X)^{-1} X^T E[Y] \tag{1}$$

$$= (X^T X)^{-1} X^T E[X\beta + \varepsilon] \tag{2}$$

$$= (X^T X)^{-1} (X^T X)\beta \tag{3}$$

$$= \beta \tag{4}$$

Suppose that $\tilde{\alpha}$ is some other linear unbiased estimator of α . Since the estimator is linear, there exists some element w such that $\tilde{\alpha} = \langle y, w \rangle$

$$\langle \beta, v \rangle = \alpha = E[\tilde{\alpha}] = E[\langle y, w \rangle] = \langle X\beta, w \rangle = \langle \beta, X^T w \rangle$$

This implies that $v = X^T w$, therefore we have,

$$\text{Var} = \text{Var}(\langle x\beta, w \rangle + \langle \varepsilon, w \rangle) \quad (5)$$

$$= \text{Var}(\langle \varepsilon, w \rangle) + E \left[\left(\sum_{i=1}^n \varepsilon_i w_i \right)^2 \right] \quad (6)$$

$$= \sigma^2 \sum_{i=1}^p w_i^2 = \sigma^2 \|w\|^2 \quad (7)$$

$$\text{Var}(\hat{\alpha}) = E[\langle \hat{\beta} - \beta, v \rangle^2] \quad (8)$$

$$= E[\langle (X^T X)^{-1} X^T \beta + (X^T X)^{-1} X^T \varepsilon - \beta, v \rangle^2] \quad (9)$$

$$= E[\langle (X^T X)^{-1} X^T \varepsilon, v \rangle^2] \quad (10)$$

$$= \sigma^2 \|X(X^T X)^{-1} v\|^2 = \sigma^2 \|X(X^T X)^{-1} X^T w\|^2 \quad (11)$$

$$= \sigma^2 \|\Pi_u w\|^2 \quad (12)$$

Thus, $\text{Var}(\hat{\alpha}) \leq \text{Var}\tilde{\alpha}$ □

7 Lecture 8

Recall linear model

$$Y = X\beta + \varepsilon$$

where $\text{cov}(\varepsilon) = \Sigma$.

OLD: $\hat{\beta} = (X_\Sigma^T X_\Sigma)^{-1} X_\Sigma^T \Sigma^{-1/2} Y$.

$X\hat{\beta}$ = Projection of $\Sigma^{-1/2} Y$ onto $\text{span} \{X_{\varepsilon,1}, \dots, X_{\varepsilon,p}\}$

Theorem 2 (Gauss-Markov). 1. $\hat{\beta}_{OLS}$ is the best linear unbiased est (BLUE)

2. $\alpha_i = \langle \beta, v \rangle$ is BLUE.

3. $\hat{\sigma}^2 = \frac{\|Y - X\hat{\beta}\|^2}{n-p}$ is unbiased est for $\sigma^2 > 0$

$$\begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{pmatrix} = \begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{pmatrix} \begin{pmatrix} \beta_1 \\ \beta_2 \\ \vdots \\ \beta_p \end{pmatrix}^T + \begin{pmatrix} \varepsilon_1 \\ \varepsilon_2 \\ \vdots \\ \varepsilon_n \end{pmatrix} \quad \text{Where our data is } (Y_i, X_i)_{i=1}^n \in (\mathbb{R} \times \mathbb{R}^p)^{\otimes p}$$

Remark 5. Is this an iid model? Depends!

1. Typically ε_i are iid.

2. If X_i are random then "random design".

3. If X_i are iid, then linear model is iid model.

4. If X_i are deterministic, then not iid model.

$$\beta \mapsto \|Y - X\hat{\beta}\|.$$

Proof. This is a continuation of point 3 in our theorem above.

We already introduced $\Pi_U = X(X^T X)^{-1} X^T$ projection onto col space U of X . Thus $I_n - \Pi_U$ is another projection operator, onto U^\perp (orthogonal complement),

$$U^\perp = \{z \in \mathbb{R}^n \mid \langle z, X_k \rangle \forall k = 1, \dots, p\}.$$

Choose a basis e_1, \dots, e_{n-p} , orthonormal, of U^\perp , then

$$(I_n - \Pi_U)z = \Pi_{U^\perp} z = \sum_{k=1}^{n-p} \langle z, e_k \rangle e_k.$$

$$\|Y - X\hat{\beta}\| = \|Y - \underbrace{X(X^T X)^{-1} X^T Y}_{\Pi_U}\| \quad (13)$$

$$= \|(I_n - \Pi_U)Y\|^2 \quad (14)$$

$$= \|(I_n - \Pi_U)(X\beta + \varepsilon)\|^2 \quad (15)$$

$$= \|(I_n - \Pi_U)\varepsilon\|^2 \quad (16)$$

$$= \sum_{i=1}^{n-p} \langle \varepsilon, e_i \rangle^2 \quad (17)$$

$$(18)$$

Hence,

$$E[\|Y - X\hat{\beta}\|^2] = \sum_{i=1}^{n-p} E[\langle \varepsilon, e_i \rangle^2] = n - p \implies E[\hat{\sigma}] = n - p$$

□

Remark 6. Recall the $N(\mu, \sigma^2)$ model, where the MLE is

$$\hat{\mu} = \frac{1}{n} \sum_{i=1}^n X_i, \quad \hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (X_i - \hat{\mu})^2.$$

The unbiased estimator for σ^2 was $\frac{1}{n-1} \sum_{i=1}^n (X_i - \hat{\mu})^2$. This is related to the $n - p$ factor in point 3.

Remark 7. 1. If linearity is dropped, there exists better estimators than $\hat{\beta}_{OLS}$. For example a constant estimator, $\hat{\beta} = \beta^*$

2. The MSE of $\hat{\beta}_{OLS}$ is

$$E[\|\hat{\beta}_{OLS} - \beta\|^2] = E\left[\sum_{i=1}^p \langle \hat{\beta}_{OLS} - \beta, \underbrace{e_i}_{\text{ONB of } \mathbb{R}^n} \rangle^2\right] = \sum_{i=1}^p \text{Var}_\beta(\langle \hat{\beta}_{OLS}, e_i \rangle) = \sum_{i=1}^p \sigma^2 \|X(X^T X)^{-1} e_i\|^2$$

We say X satisfies orthogonal design if

$$X^T X = nI_p$$

"The different covariants are uncorrelated." $(X^T X)_{ij} = \langle X_i, X_j \rangle = n\delta_{ij}$ For orthogonal design,

$$E_\beta[\|\hat{\beta}_{OLS} - \beta\|^2] = \frac{1}{n^2} \sigma^2 \sum_{i=1}^p \underbrace{\|Xe_i\|^2}_n = \frac{\sigma^2 p}{n}.$$

and this is equal to noise level times the number of parameters, divided by the number of data points.

Theorem 3 (Bayes in Linear Models). Consider a linear model $Y = X\beta + \varepsilon$, and $\varepsilon \sim N(0, \sigma^2 I_n)$ with $\sigma > 0$ known and $\beta \sim N(m, M)$ where $m \in \mathbb{R}^p, M \in \mathbb{R}^{p \times p}$ positive semi definite. Then, the posterior $\Pi(\beta|Y, X)$ is given by

$$\Pi(\beta|Y, X) = N(\mu_{past}, \Sigma_{past}) \text{ for}$$

$$\mu_{past} = \sigma_{past}^{-2} X^T y + M^{-1} m \quad \Sigma_{past} = (\sigma_{past}^{-2} X^T X + M^{-1})^{-1}$$

Remark 8. Σ_{past} independent of Y . For " $M^{-2} \rightarrow 0$ ", then " $\mu_{past} \rightarrow \hat{\beta}_{OLS}$ "

Proof.

$$L(X, Y, \beta) \pi(\beta) \propto \exp \left(-\frac{1}{2\sigma^2} \|Y - X\beta\|^2 - \frac{1}{2} (\beta - m)^T M^{-1} (\beta - m) \right)$$

We want this to be proportional to $\exp \left(-\frac{1}{2} (\beta - \mu_{\text{past}})^T \sigma_{\text{past}}^{-1} (\beta - \mu_{\text{past}}) \right)$.

Now,

$$\exp \left(-\frac{1}{2} (\beta - \mu_{\text{past}})^T \sigma_{\text{past}}^{-1} (\beta - \mu_{\text{past}}) \right) \propto \exp \left(-\frac{1}{\sigma^2} \beta^T X^T X \beta - \frac{1}{2} \beta^T M^{-1} \beta + \frac{1}{\sigma^2} \beta^T X^T Y + \beta^T M^{-1} m \right)$$

and this is equal to

$$\exp \left(-\frac{1}{2} \beta^T \left(\frac{1}{\sigma^2} X^T X + M^{-1} \right) \beta + \beta^T \left(\frac{1}{\sigma^2} X^T Y + M^{-1} m \right) \right)$$

and this is

$$\propto \exp \left(-\frac{1}{2} (\beta - \mu_{\text{past}})^T \sigma_{\text{past}}^{-1} (\beta - \mu_{\text{past}}) \right)$$

□

Corollary 1. For $\ell = \|\cdot\|^2$, the Bayes estimator is $\hat{\beta}_{\Pi} = \mu_{\text{past}}$

Proposition 1. Consider the previous setting (from the theorem), with $m = 0$, and $M = \tau^2 I_p$ (centered, isotropic, normal prior). The, $\mu_{\text{past}} = \hat{\beta}_{\Pi}$ minimizes

$$\beta \mapsto \|Y - X\beta\|_{\mathbb{R}^n}^2 + \underbrace{\frac{\sigma^2}{\tau^2} \|\beta\|_{\mathbb{R}^p}^2}_{\text{"penalty" or "regularization"}}$$

Proof.

□

8 Lecture 9

Proof:

Take gradient of $\mathcal{J}(\beta)$ w.r.t. β :

$$\nabla_{\beta} \mathcal{J}(\beta) = 2\mathbf{X}^{\top}(\mathbf{Y} - \mathbf{X}\beta) + \frac{2\sigma^2}{\tau^2} \beta$$

Set = 0:

$$\Rightarrow \nabla_{\beta} \mathcal{J}(\beta) = 2(\mathbf{X}^{\top} \mathbf{X} + \frac{\sigma^2}{\tau^2} \mathbf{I}) \beta - 2\mathbf{X}^{\top} \mathbf{Y} = 0$$

$$\Rightarrow \beta = (\mathbf{X}^{\top} \mathbf{X} + \frac{\sigma^2}{\tau^2} \mathbf{I})^{-1} \mathbf{X}^{\top} \mathbf{Y}$$

Posterior mean:

$$\begin{aligned} \mu_{\text{post}} &= \Sigma_{\text{post}}^{-1} (\mathbf{X}^{\top} \mathbf{Y} + \mathbf{M}_0^{-1} \mu_0) \\ &= (\sigma^{-2} \mathbf{X}^{\top} \mathbf{X} + \tau^{-2} \mathbf{I}_p)^{-1} \sigma^{-2} \mathbf{X}^{\top} \mathbf{Y} \\ &= (\mathbf{X}^{\top} \mathbf{X} + \frac{\sigma^2}{\tau^2} \mathbf{I})^{-1} \mathbf{X}^{\top} \mathbf{Y} \end{aligned}$$

Remark:

β is defined even if $\text{rank}(\mathbf{X}) < p$, in particular even for $n < p$.

Definition:

$$\hat{\beta}_{\text{ridge}} = \underset{\beta \in \mathbb{R}^p}{\operatorname{argmin}} \|\mathbf{Y} - \mathbf{X}\beta\|^2 + \lambda \|\beta\|^2,$$

is called a **Ridge Regression** estimator. Here, $\lambda > 0$ is called a regularization parameter. $\hat{\beta}_{\text{ridge}}$ is always uniquely defined.

For $\mathbf{Y} = \mathbf{X}\beta + \varepsilon$, UM:

$$\hat{\beta}_{\text{ridge}} = (\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I}_p)^{-1} \mathbf{X}^\top \mathbf{Y}.$$

Estimator independent of σ^2 .

Proposition:

MSE of $\hat{\beta}_{\text{ridge}}$.

Consider a linear model with $\varepsilon \sim N(0, \sigma^2 \mathbf{I}_n)$, $\sigma^2 > 0$ known, and $\mathbf{X}^\top \mathbf{X} = n \mathbf{I}_p$ (orthonormal design).

Let $\mathcal{J} := \langle \beta, \mathbf{v} \rangle$ for $\mathbf{v} \in \mathbb{R}^p$, and:

$$\delta_{\text{ridge}} = \langle \hat{\beta}_{\text{ridge}}, \mathbf{v} \rangle.$$

Then:

1.

$$\mathbb{E}_{\beta}[(\delta_{\text{ridge}} - \mathcal{J})^2] = (1 + \lambda)^{-2} \langle \beta, \mathbf{v} \rangle^2 + \frac{\sigma^2}{n} \|\mathbf{v}\|^2 (1 + \lambda)^{-2}.$$

2.

$$\mathbb{E}_{\beta}[\|\hat{\beta}_{\text{ridge}} - \beta\|^2] = (1 + \lambda)^{-2} \|\beta\|^2 + \frac{p\sigma^2}{n} \frac{1}{(1 + \lambda)^2}.$$

We have:

$$\begin{aligned} \hat{\beta}_{\text{ridge}} &= (\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I}_p)^{-1} \mathbf{X}^\top \mathbf{Y} \\ &= (n \mathbf{I}_p + \lambda \mathbf{I}_p)^{-1} \mathbf{X}^\top \mathbf{Y} \\ &= \frac{1}{(1 + \frac{\lambda}{n})} (\mathbf{X}^\top \mathbf{X} \beta + \mathbf{X}^\top \varepsilon), \end{aligned}$$

where $\varepsilon \sim N(0, \sigma^2 \mathbf{I}_n)$.

$$\begin{aligned} &= \frac{1}{1 + \frac{\lambda}{n}} (\mathbf{X}^\top \mathbf{X} \beta + \mathbf{X}^\top \varepsilon), \\ &= \frac{1}{1 + \frac{\lambda}{n}} \beta + \frac{1}{1 + \frac{\lambda}{n}} \mathbf{X}^\top \varepsilon. \end{aligned}$$

Bias-Variance Decomposition:

$$\begin{aligned} \mathbb{E}[(\hat{\beta}_{\text{ridge}} - \mathcal{J})^2] &= (\mathbb{E}[\hat{\beta}_{\text{ridge}}] - \mathcal{J})^2 + \operatorname{Var}(\hat{\beta}_{\text{ridge}}). \\ &= ((1 + \frac{\lambda}{n})^{-1} \langle \beta, \mathbf{v} \rangle)^2 + \frac{\lambda^2}{(1 + \lambda)^2} \operatorname{Var}(\mathbf{X}^\top \varepsilon, \mathbf{v}). \end{aligned}$$

Observe:

$$(1 + \frac{\lambda}{n})^{-1} = \frac{1}{(1 + \frac{\lambda}{n})}.$$

Also:

$$\operatorname{Var}(\mathbf{X}^\top \varepsilon, \mathbf{v}) = \mathbf{v}^\top \mathbf{X} \operatorname{Cov}(\varepsilon) \mathbf{X}^\top \mathbf{v} = \sigma^2 \|\mathbf{v}\|^2.$$

Corollary:

Under the same assumptions:

$$\begin{aligned} \mathbb{E}[\|\hat{\beta}_{\text{ridge}} - \beta\|^2] &= \mathbb{E} \left[\sum_{k=1}^p (\langle \beta, \mathbf{e}_k \rangle - \beta_k)^2 \right] \\ &= \frac{1}{(1 + \frac{\lambda}{n})^2} \|\beta\|^2 + \frac{p\sigma^2}{n(1 + \frac{\lambda}{n})^2}. \end{aligned}$$

Remark:

For small $\|\beta\|$, Ridge \rightarrow OLS. The optimal choice of λ depends on $\|\beta\|$.

1.7 Confidence Sets & Tests in Linear Model:

The estimators we studied are independent of σ^2 , but uncertainty quantification will depend on σ^2 !

Assume $\varepsilon \sim N(0, \sigma^2 \mathbf{I}_n)$ throughout.

Easy Case:

For $\sigma^2 > 0$ known:

$$\hat{\beta}_{\text{OLS}} \sim N(\beta, (\mathbf{X}^\top \mathbf{X})^{-1}).$$

Indeed:

$$\text{Cov}((\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \varepsilon) = (\mathbf{X}^\top \mathbf{X})^{-1}.$$

And for $\mathcal{J} = \langle \beta, \nu \rangle$,

$$\hat{\mathcal{J}} = \langle \hat{\beta}_{\text{OLS}}, \nu \rangle \sim N(\mathcal{J}, \sigma^2 \nu^\top (\mathbf{X}^\top \mathbf{X})^{-1} \nu).$$

Then a 95% confidence set for \mathcal{J} is:

$$I_{95\%}(\mathcal{J}) = \left[\hat{\mathcal{J}} \pm 1.96 \sqrt{\nu^\top (\mathbf{X}^\top \mathbf{X})^{-1} \nu} \right].$$

Notes on t - and F -distributions:

BUT: Normally, σ^2 is unknown. Replace σ by its estimator $\hat{\sigma}$. We need the t - and F -distributions.

Definitions:

Definition (t-distribution): The t -distribution with $n \geq 1$ degrees of freedom on \mathbb{R} has density:

$$f_n(x) = C_n \left(1 + \frac{x^2}{n} \right)^{-\frac{n+1}{2}},$$

where C_n is the normalizing constant.

Note: For $n = 1$:

$$f_1(x) = C_1 \frac{1}{1 + x^2},$$

which corresponds to the **Cauchy distribution**.

Definition (F-distribution): The F -distribution with $(m, n) \in \mathbb{N}^2$ degrees of freedom has density:

$$f_{m,n}(x) = C_{m,n} \frac{x^{\frac{m}{2}-1}}{(mx + n)^{\frac{m+n}{2}}}, \quad x \in (0, \infty),$$

where $C_{m,n}$ is the normalizing constant.

Why is this useful?

Lemma: Let $X_1, \dots, X_n, Y_1, \dots, Y_n$ be i.i.d. $N(0, \Delta)$ random variables. Then:

1.

$$T_n := \frac{X_n}{\sqrt{\frac{1}{n} \sum_{i=1}^n Y_i^2}} \sim t_n.$$

2.

$$F_{m,n} := \frac{\frac{1}{m} \sum_{i=1}^m X_i^2}{\frac{1}{n} \sum_{j=1}^n Y_j^2} \sim F_{m,n}.$$

Remarks:

1. The t -distribution arises when considering the "empirical mean" and "empirical variance."
2. For $n \rightarrow \infty$, $T_n \xrightarrow{d} N(0, 1)$.

Proof:**(b) Observe:**

$$T_n^2 = F_{1,n}.$$

By a change of measure ($y \mapsto y^2$ in $(0, \infty)$):

$$f_{F_{m,n}}(x) = f_{F_{m,n}}(x^2)2x, \quad x > 0.$$

Since t is symmetric around 0, we obtain for all $x \in \mathbb{R}$:

$$f_{T_n}(x) = f_{F_{m,n}}(x^2)|x| = C_n \left(1 + \frac{x^2}{n}\right)^{-\frac{n+1}{2}}.$$

It remains to show the claim for $F_{m,n}$.

Let:

$$X = \sum_{i=1}^m X_i^2, \quad Y = \sum_{j=1}^n Y_j^2.$$

Then:

$$X \sim \chi_m^2, \quad Y \sim \chi_n^2,$$

where the density of χ_m^2 is:

$$f(x) \propto x^{m/2-1} e^{-x/2}, \quad x > 0.$$

Derivation:

Writing $W = \frac{X}{Y}$, we have:

$$\mathbb{P}\left(\frac{X}{Y} < z\right) = \int_0^\infty \int_0^{zy} 1 f_X(x) f_Y(y) dx dy.$$

Substituting $x = wy$, we get:

$$\begin{aligned} &= \int_0^\infty \int_0^z 1 f_X(wy) f_Y(y) y dw dy \\ &= \int_0^\infty f_X(zy) f_Y(y) y dy \\ &\propto \int_0^\infty (zy)^{\frac{m}{2}-1} y^{\frac{n}{2}-1} e^{-(z+y)/2} dy. \end{aligned}$$

Change of Variable:

Let $a = \frac{z}{z+1}y$, then:

$$\begin{aligned} &\propto \int_0^\infty \left(\frac{z}{z+1}\right)^{\frac{m}{2}} a^{\frac{m}{2}-1} e^{-\frac{z}{z+1}a} \frac{1}{z+1} da \\ &\propto z^{\frac{m}{2}-1} (z+1)^{-\frac{m+n}{2}} \int_0^\infty a^{\frac{m}{2}-1} e^{-a} da. \end{aligned}$$

It follows:

$$\begin{aligned} \frac{\partial}{\partial z} \mathbb{P}\left(\frac{X}{Y} < z\right) &= f_{X,Y}(z) = \int_0^\infty f_X(zy) f_Y(y) \frac{1}{y} dy \\ &\propto z^{\frac{m}{2}-1} (z+1)^{-\frac{m+n}{2}}. \end{aligned}$$

Change of Variable:

Let $F = \frac{X}{Y}$, given $f_F(z) = \frac{m}{n} f_{X,Y} \left(\frac{m}{n} z \right) = f_{m,n}(z)$.

9 Lecture 10

t -distribution $\cdot t_n(x) \propto \left(\frac{n^2}{n} + 1 \right)^{-(n+1)/2}$

$$F \cdot f_{m,n}(x) \propto \frac{x^{m/2-1}}{(n+mx)^{(m+n)/2}}$$

In the linear model $Y = X\beta + \varepsilon$, $\varepsilon \sim N(0, \sigma^2 I_n)$,

$$\hat{\beta} = (X^T X)^{-1} X^T Y \sim N(0, \sigma^2 A)$$

where

$$\sigma^2 = \hat{\sigma}^2 = \left\| \frac{Y - X\hat{\beta}}{n-p} \right\|^2 \sim \sigma^2 \frac{\chi^2(n-p)}{n-p}$$

We now have

$$t_n = \frac{N(0,1)}{\chi^2(n)n} \quad f_{m,n} = \frac{n\chi^2(m)}{m\chi^2(n)}$$

Lemma 2. Let $\xi \sim N(0, I_n)$, a random variable in \mathbb{R}^n , and let $R \in \mathbb{R}^{n \times n}$ be an orthogonal projection ($R = R^2, R = R^T$), with $\text{rank}(R) = r \leq n$.

1. $\xi^T R \xi = \|R\xi\|^2 \sim \chi^2(r)$.
2. If $B \in \mathbb{R}^{p \times n}$ is such that $BR = 0$, then $B\xi$ is independent from $R\xi$
3. If $S \in \mathbb{R}^{n \times n}$ is another orthogonal projection, $\text{rank}(S) = s \leq n$ and $RS = 0$, then

$$\frac{s}{r} \frac{\xi^T R \xi}{\xi^T S \xi} \sim F(r, s)$$

Proof. 1. Since R is an orthogonal projection, there exists an orthogonal matrix $T^T = T^{-1}$ such that

$$R = T \begin{pmatrix} I_r & \mathbf{0} \\ \mathbf{0} & \mathbf{0} \end{pmatrix} T^T = T D_r T^T.$$

Then we have $T^T \sim N(0, T^T T) = N(0, I_n)$.

$$\xi^T R \xi = \xi^T (T D_r T^T) \xi = (T^T \xi)^T D_r T^T \xi = \sum_{i=1}^n (T^T \xi)_i^2 \sim \chi^2(r).$$

2. Let $A_1 = B\xi$, $A_2 = R\xi$, then

$$\text{Cov}(A_1, A_2) = \text{Cov}(B\xi, R\xi) = B \text{Cov}(\xi, \xi) R^T = B R^T = B R = 0$$

3. By (2), we know $S\xi$ and $R\xi$ are independent. By (1), $\xi^T S \xi \sim \chi^2(s)$, $\xi^T R \xi \sim \chi^2(r)$. The claim follows from the definition of $F(r, s)$. □

Theorem 4. Linear Model Confidence Sets -unknown σ^2 Assume regular linear model, $Y = X\beta + \varepsilon$, $\text{rank}(X) = p \leq n$, $\varepsilon \sim N(0, \sigma^2 I_n)$. Let $\alpha \in (0, 1)$

1. Let $q_{F_{p,n-p}, 1-\alpha}$ be the $1-\alpha$ quantile of $F_{p,n-p}$ distribution. Then $C(Y, X) = \left\{ \beta \in \mathbb{R}^p \mid \frac{\|X(\beta - \hat{\beta}_{OLS})\|^2}{p\hat{\sigma}^2} \leq q_{F_{p,n-p}, 1-\alpha} \right\}$ is a level $1-\alpha$ confidence set.
2. Let $\alpha = \langle \beta, v \rangle$, for some $v \in \mathbb{R}^p$. Then a $1-\alpha$ confidence set is

$$C = C(Y, X) = \left\{ \alpha \in \mathbb{R} \mid \left| \frac{\alpha - \hat{\alpha}}{\hat{\sigma} \sqrt{v^T (X^T X)^{-1} v}} \right| < q \right\}$$

where $\hat{\alpha} = \langle \hat{\beta}_{OLS}, v \rangle$ and q is the $1-\alpha/2$ quantile of t_n .

Proof. 1. We know $X\hat{\beta}_{OLS} = \Pi_U Y = \Pi_U X\beta + \Pi_U \varepsilon = X\beta + \Pi_U \varepsilon$ Moreover,

$$\hat{\sigma}^2 = \frac{\|X(\beta - \hat{\beta}_{OLS})\|^2}{n-p} = \frac{\|(I_n - \Pi_U)Y\|^2}{n-p} = \frac{\|\Pi_{U^\perp} Y\|^2}{n-p} = \frac{\|\Pi_{U^\perp} \varepsilon\|^2}{n-p}$$

This implies

$$\frac{\|X\beta - X\hat{\beta}_{OLS}\|^2}{p\hat{\sigma}^2} = \frac{(n-p)\|\Pi_U \varepsilon\|^2}{p\|\Pi_{U^\perp} \varepsilon\|^2} \sim \frac{(n-p)\sigma^2\chi^2(p)}{p\sigma^2\chi^2(n-p)} \sim F(p, n-p).$$

2. We know

$$\hat{\sigma} = \langle \hat{\beta}_{OLS}, v \rangle = v^T \hat{\beta}_{OLS} \sim v^T N(\beta, (X^T X)^{-1} \sigma^2) = N(\alpha, v^T (X^T X)^{-1} v \sigma^2)$$

And this implies

$$\frac{\alpha - \hat{\alpha}}{\sigma \sqrt{v^T (X^T X)^{-1} v}} \sim N(0, 1).$$

Finally, also, as in (1), $\hat{\sigma}^2 \sim \sigma^2 \chi^2(n-p)$. This implies

$$\frac{\alpha - \hat{\alpha}}{\hat{\sigma} \sqrt{v^T (X^T X)^{-1} v}} \sim t_{n-p}.$$

□

9.1 The t - and F -test

Remark 9 (Method (t-test)). In a regular linear model with $\varepsilon \sim N(0, \sigma I_n)$, consider $H_0 : \gamma = \gamma_0$ vs $H_1 : \gamma \neq \gamma_0$ ($\gamma = \langle \beta, v \rangle$). The two sided t -test is

$$\varphi_{\alpha_0}(Y, X) = \mathbf{1}(\{|T_{\alpha_0, n-p}(Y, X)| > q\}),$$

where

$$T_{\alpha_0, n-p} = \frac{\alpha_0 - \hat{\alpha}}{\hat{\sigma} \sqrt{v^T (X^T X)^{-1} v}}$$

and q is the $1 - \alpha/2$ -quantile of t_{n-p} .

Remark 10 (Method (F-test)). Same setting as before for t -test, $H_0 : \beta = \beta_0$ vs $H_1 : \beta \neq \beta_0$ since $\beta_0 \in \mathbb{R}^p$. Then the F -test is

$$\varphi_{\beta_0}(Y, X) = \mathbf{1}(|F_{\beta_0, n-p}(Y, X)| > q)$$

where

$$F_{\beta_0, n-p}(Y, X) = \frac{\|X(\beta - \hat{\beta}_{OLS})\|^2}{p\hat{\sigma}^2}$$

and $q = (1 - \alpha)$ -quantile of $F_{p, n-p}$.

9.2 General linear hypothesis testing problems

Definition 9. A linear hypothesis testing problem is of the form $H_0 : K\beta = d$ vs $H_1 : K\beta \neq d$, where $K \in \mathbb{R}^{r \times p}$ with $\text{rank}(K) = r \leq p$, $d \in \mathbb{R}^r$. In other words "r linear constraints on β "

K is called the "contrast matrix"

Theorem 5. Assume regular linear model, with $\varepsilon \sim N(0, \sigma^2 I_n)$, and consider $H_0 : K\beta = d$ vs. $K\beta \neq d$.

Define residual sum of squares as $RSS = \|Y - X\hat{\beta}_{OLS}\|^2$ and $RSS_{H_0} = \|Y - X\hat{\beta}_{H_0}\|^2$ and $\hat{\beta}_{H_0}$ over $\{\beta : K\beta = d\}$.

Proof. 1.

$$\hat{\beta}_{H_0} = \hat{\beta}_{OLS} - (X^T X)^{-1} K^T (K(X^T X)^{-1} K^T)^{-1} (K\hat{\beta}_{OLS} - d)$$

$$2. RSS_{H_0} - RSS = (K\hat{\beta}_{OLS} - d)(K(X^T X)^{-1} K^T)^{-1} (K\hat{\beta}_{OLS} - d), \quad \frac{RSS_{H_0} - RSS}{\sigma^2} \sim \chi^2(r)$$

3. Define

$$F = \frac{n-p}{r} = \frac{RSS_{H_0} - RSS}{RSS} = \frac{RSS_{H_0} - RSS}{r\hat{\sigma}^2} \sim F_{r, n-p}$$

under H_0 .

□

10 Lecture 11

Theorem:

Assume regular LM, $\varepsilon \sim N(0, \sigma^2 I_n)$, and consider:

$$H_0 : K\beta = d \quad \text{vs.} \quad H_1 : K\beta \neq d$$

Define:

$$RSS = \|Y - X\beta\|^2, \quad RSS_{H_0} = \|Y - X\beta_{H_0}\|^2$$

where β_{H_0} is the OLS estimator over $K\beta = d$:

$$\beta_{H_0} = \hat{\beta} - (X^\top X)^{-1} K^\top (K(X^\top X)^{-1} K^\top)^{-1} (K\hat{\beta} - d)$$

1.

$$RSS_{H_0} - RSS = \|X(\beta_{H_0} - \hat{\beta})\|^2 = (K\hat{\beta} - d)^\top (K(X^\top X)^{-1} K^\top)^{-1} (K\hat{\beta} - d)$$

2. Under H_0 :

$$RSS_{H_0} \sim \chi^2(n)$$

3. Define:

$$F = \frac{1}{p} \frac{RSS_{H_0} - RSS}{RSS/n} = \frac{RSS_{H_0} - RSS}{c \cdot RSS}$$

Under H_0 :

$$F \sim F_{n-p}$$

Proof:

1. To show $K\hat{\beta}_{H_0} = d$, note that β_{H_0} is the minimizer.

Observe:

$$K\beta_{H_0} - K\beta = K(X^\top X)^{-1} K^\top (K(X^\top X)^{-1} K^\top)^{-1} (K\beta - d) - K\beta - (K\beta - d) = d$$

Second part: Let $Y \in \mathbb{R}^n$, $K\beta = d$. By Pythagoras:

$$\|Y - X\hat{\beta}\|^2 = \|Y - X\beta_{H_0}\|^2 + \|X(\hat{\beta} - \beta_{H_0})\|^2$$

where:

$$A = \left(X(\hat{\beta} - \beta_{H_0}) \right)^\top X\beta_{H_0} - Y = (K\hat{\beta} - d)^\top (K(X^\top X)^{-1} K^\top)^{-1} K(X^\top X)^{-1} (X^\top Y)$$

This implies:

$$(K\hat{\beta} - d)^\top (K(X^\top X)^{-1} K^\top)^{-1} (K\hat{\beta} - d) = 0$$

Overall:

$$\|Y - X\beta_{H_0}\|^2 = \|Y - X\hat{\beta}\|^2 + \|X(\hat{\beta} - \beta_{H_0})\|^2 \geq 0$$

Continuation:

2) Under H_0 :

$$\begin{aligned} RSS_{H_0} - RSS &= \|Y - X\beta_{H_0}\|^2 - \|Y - X\hat{\beta}\|^2 \\ &= \|X(\hat{\beta} - \beta_{H_0})\|^2 = (K\hat{\beta} - d)^\top (K(X^\top X)^{-1} K^\top)^{-1} (K\hat{\beta} - d) \end{aligned}$$

Let $Z = K\hat{\beta} - d$. Then:

$$\mathbb{E}[Z] = \mathbb{E}[K\hat{\beta} - d] = K\mathbb{E}[\hat{\beta}] - d = K(X^\top X)^{-1} X^\top Y - d$$

(Substitute $K\beta = d$ into the expectation)

$$\text{Var}(Z) = K\text{Var}(\hat{\beta})K^\top = \sigma^2 K(X^\top X)^{-1} K^\top$$

Thus:

$$Z \sim \mathcal{N}(0, \sigma^2 K(X^\top X)^{-1} K^\top)$$

Finally:

$$RSS_{H_0} - RSS = \|X(\hat{\beta} - \beta_{H_0})\|^2 = Z^\top (\sigma^2 K(X^\top X)^{-1} K^\top)^{-1} Z$$

$$\sim \chi^2(p)$$

$$RSS \sim \sigma^2 \chi^2(n-p), \quad RSS_{H_0} \sim \sigma^2 \chi^2(n).$$

3) We know:

$$\frac{RSS_{H_0} - RSS}{\sigma^2} \sim \chi^2(p), \quad \frac{RSS}{\sigma^2} \sim \chi^2(n-p)$$

To show independence: We have:

$$RSS_{H_0} \perp Y \quad \text{while} \quad RSS_{H_0} - RSS \text{ only depends on } \hat{\beta}.$$

(Since $\hat{\beta} \propto X^\top Y$ and $T_n = 0$ by the lemma from last time, independence follows.)

ANOVA (Analysis of Variance)

Motivation: We have data from k different groups. Are the means equal?

Definition: ANOVA

We are given data:

$$Y_{ij} = \mu_i + \varepsilon_{ij}, \quad i = 1, \dots, k, \quad j = 1, \dots, n_i$$

Assume:

$$\varepsilon_{ij} \sim \mathcal{N}(0, \sigma^2)$$

ANOVA (Analysis of Variance)

Index: $i = 1, \dots, k$ is called the factor.

The model is a *factor model* with 1 categorical variable.

$n = \sum_{i=1}^k n_i$ is the total sample size.

The model is balanced (design) if $n_1 = n_2 = \dots = n_k$.

Remark: ANOVA is a linear model:

$$\begin{pmatrix} Y_{1,1} \\ Y_{1,2} \\ \vdots \\ Y_{k,n_k} \end{pmatrix} = \begin{pmatrix} \mathbf{1}_{n_1} & 0 \\ 0 & \mathbf{1}_{n_k} \end{pmatrix} \begin{pmatrix} \mu_1 \\ \mu_k \end{pmatrix} + \varepsilon$$

Hypothesis Testing:

$$H_0 : \mu_1 = \dots = \mu_k \quad \text{vs.} \quad H_a : \exists i, j \text{ with } \mu_i \neq \mu_j$$

Basic Idea: Compare variation within groups vs. variation across groups.

Theorem (Decomposition of RSS): Define the group means:

$$\bar{Y}_{i.} = \frac{1}{n_i} \sum_{j=1}^{n_i} Y_{ij}, \quad i = 1, \dots, k$$

and the overall mean:

$$\bar{Y}_{..} = \frac{1}{n} \sum_{i,j} Y_{ij}.$$

Furthermore, let:

$$SSB = \sum_{i=1}^k n_i (\bar{Y}_{i.} - \bar{Y}_{..})^2 \quad (\text{Sum of squares between groups}),$$

$$SSW = \sum_{i=1}^k \sum_{j=1}^{n_i} (Y_{ij} - \bar{Y}_{i.})^2 \quad (\text{Sum of squares within groups}).$$

Then:

$$SST = \sum_{i=1}^k \sum_{j=1}^{n_i} (Y_{ij} - \bar{Y}_{..})^2 = SSB + SSW,$$

where SST is the total sum of squares.

Proof:

$$\begin{aligned} SST &= \sum_{i=1}^k \sum_{j=1}^{n_i} (Y_{ij} - \bar{Y}_{i.} + \bar{Y}_{i.} - \bar{Y}_{..})^2 \\ &= \sum_{i,j} (Y_{ij} - \bar{Y}_{i.})^2 + (\bar{Y}_{i.} - \bar{Y}_{..})^2 + 2(Y_{ij} - \bar{Y}_{i.})(\bar{Y}_{i.} - \bar{Y}_{..}) \\ &= SSB + SSW + C, \end{aligned}$$

where:

$$C = \sum_{i=1}^k (\bar{Y}_{i.} - \bar{Y}_{..}) \sum_{j=1}^{n_i} (Y_{ij} - \bar{Y}_{i.}).$$

By construction:

$$\sum_{j=1}^{n_i} (Y_{ij} - \bar{Y}_{i.}) = 0 \quad \text{for each } i,$$

so:

$$C = 0.$$

Theorem:

1. The least square estimator for $\mu = (\mu_1, \dots, \mu_k) \in \mathbb{R}^k$ is:

$$\hat{\mu} = (\bar{Y}_{1.}, \dots, \bar{Y}_{k.})^\top.$$

2. Under H_0 :

$$\frac{SSW}{\sigma^2} \sim \chi^2(n - k).$$

3. Under H_0 :

$$\frac{SSB}{\sigma^2} \sim \chi^2(k - 1).$$

4. SSW and SSB are independent under H_0 , and:

$$F = \frac{\frac{n-k}{k-1} SSB}{SSW} \stackrel{H_0}{\sim} F(k - 1, n - k).$$

11 Lecture 12

ANOVA

linear model, factor/category, F -test for equality of means, $Y_{i,j} = \mu_i + \varepsilon_{ij}$ for $i = 1, \dots, k$ and $j = 1, \dots, n_i$.

First a note, $X^T X = \|X\|_{\mathbb{R}^n}^2$

Theorem 6. In the ANOVA model with $\varepsilon_{ij} \sim N(0, \sigma^2)$:

1. The OLS estimate is

$$\hat{\mu} = (\bar{y}_{1.}, \dots, \bar{y}_{k.}) \quad \text{Recall } \bar{y}_{i.} = \frac{1}{n_i} \sum_{j=1}^{n_i} y_{ij}$$

2.

$$\frac{SSW}{\sigma^2} = \frac{1}{\sigma^2} \sum_i \sum_j (y_{ij} - \bar{y}_{i.})^2 \sim \chi^2(n - k)$$

3. Under H_0 : $\mu_0 = \mu_1 = \dots = \mu_k$, $\frac{SSB}{\sigma^2} = \frac{1}{\sigma^2} \sum_i n_i (\bar{y}_{i\cdot} - \bar{y}) \sim \chi^2(k-1)$

4. SSW and SSB are independent and under H_0 ,

$$\frac{n-k}{k-1} \frac{SSB}{SSW} \sim F(k-1, n-k)$$

Proof. (a) We have $\hat{\mu} = (X^T X)^{-1} X^T Y$, with

$$(X^T X)^{-1} = \begin{pmatrix} \frac{1}{n_1} & \dots & 0 \\ 0 & \dots & 0 \\ 0 & \dots & \frac{1}{n_k} \end{pmatrix}$$

and this implies that

$$\hat{\mu} = \begin{pmatrix} \frac{1}{n_1} & \dots & 0 \\ 0 & \dots & 0 \\ 0 & \dots & \frac{1}{n_k} \end{pmatrix} \begin{pmatrix} \mathbb{1} & \dots & 0 \\ 0 & \dots & 0 \\ 0 & \dots & \mathbb{1} \end{pmatrix} Y = \begin{pmatrix} \frac{1}{n_1} & \dots & 0 \\ 0 & \dots & 0 \\ 0 & \dots & \frac{1}{n_k} \end{pmatrix} \begin{pmatrix} \sum_j Y_{1,j} \\ \vdots \\ \sum_j Y_{k,j} \end{pmatrix} = \begin{pmatrix} Y_{1,\cdot} \\ \vdots \\ Y_{k,\cdot} \end{pmatrix}$$

(b)

$$SSW = \|Y - X\hat{\mu}\|_{\mathbb{R}^n}^2 = \sum_i \sum_j (y_j - y_{i,\cdot})^2 = RSS \implies \frac{SSW}{\sigma^2} \sim \chi^2(n-k).$$

(c) We know that $SSW = RSS$ and we know that $SSW + SSB = SST = \sum_{ij} (y_{ij} - \bar{y}_{i\cdot})^2$.

We also know $\frac{RSS_{H_0} - RSS}{\sigma^2} \sim \chi^2(k-1)$ from before, it suffices to show $SST = RSS_{H_0}$,

$$RSS_{H_0} = \min_{\mu \in \mathbb{R}} \|Y - \mu\|_{\mathbb{R}^n}^2 = \|Y - \bar{Y}_{\cdot}\|_{\mathbb{R}^n}^2 = SST.$$

(d) Follows from general lin hypotheses testing theorem, Theorem 2.2.30 in Methoden der Statistik book. □

11.1 Exponential Families

General Model($P_\theta : \theta \in \Theta$) \supseteq Exp. families \supseteq Linear Model

Regularity Assumptions:

Let $(P_\theta : \theta \in \Theta)$ be a statistical model

1. Dominated, there exists μ such that $P_\theta \ll \mu$ for all $\theta \in \Theta$
2. $\Theta \in \mathbb{R}^p$ is an open set $p \geq 1$.
3. Likelihood $p_\theta(x) > 0$ for all $\theta \in \Theta, x \in X$, in particular $\log p_\theta(x)$ is well defined.

Definition 10. *Score* The score vector is $U_\theta(x) = \nabla_\theta \log p_\theta(x) = \begin{pmatrix} \frac{\partial}{\partial \theta_1} \log p_\theta(x) \\ \vdots \\ \frac{\partial}{\partial \theta_p} \log p_\theta(x) \end{pmatrix}$ whenever it exists

Definition 11. *Fisher Information* For $\theta \in \Theta$, the FI, whenever it exists, is $I(\theta) = E(U_\theta(x) U_\theta(x)^T) \in \mathbb{R}^{p \times p}$

More Regularity Assumptions

1. $p_\theta(x)$ is twice differentiable, in particular, $U_\theta(x)$ is well defined.
2. $E_\theta[\|U_\theta(x)\|_{\mathbb{R}^p}^2] < \infty$ for all $\theta \in \Theta$, so $I(\theta)$ is well defined.
- 3.

$$\int h(x) \nabla_\theta p_\theta(x) \mu(dx) = \nabla_\theta \int h(x)$$

for relevant $h(x)$.

Lemma 3. ($P_\theta \in \Theta$) regular model (as above),

$$1. E_\theta(U_\theta(x)) = 0$$

$$2. I(\theta) = \text{Cov}(U_\theta)$$

Proof.

$$E_\theta[U_\theta(x)] = \int_X \frac{\nabla_\theta p_\theta(x)}{p_\theta(x)} p_\theta(x) d\mu(x) = \int_X \nabla_\theta p_\theta(x) d\mu(x) = \nabla_\theta[\theta \mapsto 1] = 0$$

□

Definition 12. *Uniform minimum variance unbiased estimators (UMVUE)* Let $p(\theta) \in \mathbb{R}$ be some quantity of interest, $T: X \rightarrow \mathbb{R}$ is a UMVUE if $E_\theta[T(x)] = g(\theta)$ and for all other unbiased estimators, $S: X \rightarrow \mathbb{R}$, $\text{Var}_\theta(T) \leq \text{Var}_\theta(S)$ for all $\theta \in \Theta$.

Remark 11. 1. UMVUE are the best possible among unbiased estimators.

2. Compare to Gauss Markov, $\hat{\beta}_{OLS}$ is UMVUE

3.

$$E_\theta[||T(x) - \rho(\theta)||^2] = \text{Bias}^2 + \text{Var}(T) = \text{Var}(T) \leq E_\theta[||S(x) - \rho(\theta)||^2]$$

Theorem 7. Let $(P_\theta \mid \theta \in \Theta)$ be regular. Let $\rho: \Theta \rightarrow \mathbb{R}$ continuous differentiable, then for any unbiased estimator T of S , $E_\theta[T] = S(\theta)$,

$$\text{Var}_\theta \geq \nabla_\theta \rho(\theta)^T I(\theta)^{-1} \nabla_\theta \rho(\theta)$$

Remark 12. If $\Theta = \mathbb{R}$, $\rho(\theta) = \theta$, $\text{Var}_\theta(T) \geq I(\theta)^{-1}$

Proof. Let us assume $\Theta \subseteq \mathbb{R}$,

$$\text{Cov}_\theta(U_\theta, T) = E_\theta[U_\theta T] - E_\theta[U_\theta]E_\theta[T] = E_\theta[U_\theta T]$$

More over, by Cauchy Swartz,

$$\text{Cov}_\theta(U_\theta T) \leq \text{Var}_\theta(U_\theta)^{1/2} \text{Var}_\theta(T)^{1/2} I(\theta)^{1/2} \text{Var}_\theta(T)^{1/2}$$

But then

$$\begin{aligned} E_\theta[U_\theta T] &= \int_X \nabla_\theta \log p_\theta(x) T(x) p_\theta(x) d\mu(x) \\ &= \int_X T(x) \nabla_\theta p_\theta(x) d\mu(x) \\ &= \nabla_\theta \int \int_X T(x) p_\theta(x) d\mu(x) \\ &= E_\theta[T] = \rho'(\theta) \end{aligned}$$

Thus, $\text{Var}_\theta(T) \geq I(\theta)^{-1} \rho'(\theta)^2$

□

Another regularity condition, $I(\theta)$ is invertible.

Lecture 13

Regular Stat Model

- $\Theta \subset \mathbb{R}^p$ open
- $p_\vartheta(x) > 0$ for all $\vartheta \in \Theta$, $x \in \mathcal{X}$ and p_ϑ is continuously differentiable.

$$I(\vartheta) = \mathbb{E}_\vartheta [\nabla_\vartheta \log p_\vartheta(x) \nabla_\vartheta \log p_\vartheta(x)^T]$$

exists $\forall \vartheta \in \Theta$, and $I(\vartheta)$ is positive definite ($\Rightarrow I(\vartheta)^{-1}$ exists).

Interchange ∇_ϑ and \int .

Theorem

($p_\vartheta, \vartheta \in \Theta$ regular.) Let $g : \Theta \rightarrow \mathbb{R}$ be continuously differentiable.

Let $T : \mathcal{X} \rightarrow \mathbb{R}$ be an unbiased estimator, $\mathbb{E}_\vartheta[T] = g(\vartheta) \forall \vartheta \in \Theta$.

Then

$$\text{Var}_\vartheta(T) \geq (g'(\vartheta))^T I(\vartheta)^{-1} g'(\vartheta) \quad \forall \vartheta \in \Theta.$$

Cramér-Rao / Information Inequality

Score Vector

$$U_\vartheta(x) = \nabla_\vartheta \log p_\vartheta(x)$$

Fisher Information Matrix

Remarks

- If $I(\vartheta)$ is large, better estimation seems possible: “more information contained in the data.”
- Another interpretation.

Derivation

Let $\Theta \subseteq \mathbb{R}$. Suppose p_ϑ is twice differentiable in ϑ :

$$(\log p_\vartheta(x))' = \frac{p'_\vartheta(x)}{p_\vartheta(x)}$$

$$(\log p_\vartheta(x))'' = \frac{p''_\vartheta(x)p_\vartheta(x) - (p'_\vartheta(x))^2}{p_\vartheta(x)^2}$$

$$\mathbb{E}_\vartheta [(\log p_\vartheta(x))'] = \int_x \frac{p'_\vartheta(x)}{p_\vartheta(x)} p_\vartheta(x) dx = \int_x p'_\vartheta(x) dx = \frac{d}{d\vartheta} \int_x p_\vartheta(x) dx = 0.$$

Thus,

$$\mathbb{E}_\vartheta [(\log p_\vartheta(x))^2] = -\mathbb{E}_\vartheta [(\log p_\vartheta(x))''] = -\mathbb{E}_\vartheta [U_\vartheta(x)^2] = -I(\vartheta).$$

Theorem 8. Let $(P_\theta, \theta \in \Theta)$ be a regular model, $\Theta \subseteq \mathbb{R}$ and let $\rho : \Theta \rightarrow \mathbb{R}$, be a continuous differentiable function, an unbiased estimator T , $E_\theta[T] = \rho(\theta)$ attains equality in the CR-bound iff and only if

$$T(x) = \rho(\theta) + \rho'(\theta)I(\theta)^{-1}U_\theta(x)$$

almost surely for all $\theta \in \Theta$

Proof. Define $v(\theta) = \rho'(\theta)I(\theta)^{-1}$, then let T as above,

$$\begin{aligned} 0 \leq \text{var}(T - v(\theta)U_\theta) &= \text{var}(T) + v(\theta)^2 E_\theta[U_\theta^2] - 2v(\theta) \underbrace{\text{Cov}_\theta(T, U_\theta)}_{\rho'(\theta)} \\ &= \text{Var} - \rho'(\theta)^2 I(\theta)^{-1} = 0 \end{aligned}$$

This implies

$$T - v(\theta)U_\theta = \text{Constant}$$

Since T is unbiased we have $E_\theta[T] = \rho(\theta)$ so, $T = \rho(\theta) + v(\theta)U_\theta$ almost surely. This shows \implies , \Leftarrow is a straightforward computation. \square

Remark 13. 1. $T(x)$ is not always a measurable feature of x in the equation above.

2. If T attains the CR-bound, we say that T is the Cramer-Rao coefficient

Corollary 2. Assume previous scaling and assume $\rho(\theta) \neq 0$ for all $\theta \in \Theta$ then the likelihood can be written in the form

$$p_\theta(x) = c(x) \exp(n(\theta)T(x) - \Psi(\theta))$$

where $n : \Theta \rightarrow \mathbb{R}$, such that $n'(\theta) = \frac{I(\theta)}{\rho'(\theta)}$ and $c(x)$ and $\Psi(\theta)$ are invertible.

Proof. By the above equation, from the last theorem, we have

$$T(x) = \rho(\theta) + \rho(\theta)I^{-1}(\theta)(\log p_\theta(x))'$$

and this implies

$$(T(x) - \rho(\theta)) \frac{I(\theta)}{\rho(\theta)} = (\log p_\theta(x))'$$

and then we get

$$T(x) \int_{\theta_0}^{\theta} \frac{I(\theta)}{\rho(\theta)} dt + \Psi(\theta) = \log(p_\theta(\theta)) + \text{constant}$$

which implies

$$p_\theta(x) = \exp(\text{constant}(x)) = \exp(n(\theta)T(x) - \Psi(\theta))$$

□

Definition 13. Exponential Families A regular model $(P_\theta : \theta \in \Theta)$ is called the k -parameter Exponential family ($k \geq 1$) if there exists measurable functions

1. $n : \Theta \rightarrow \mathbb{R}^k$
2. $T : C \rightarrow \mathbb{R}^k$
3. $c : X \rightarrow [0, \infty)$

such that

$$p_\theta(x) = \frac{dP_\theta}{d\mu}(x) = c(x) \exp(\langle n(\theta)T(x) \rangle_{\mathbb{R}^k} - \Psi(\theta))$$

for all θ, x where

$$\Psi(\theta) = \log \left(\int_X c(x) \exp(\langle n(\theta)T(x) \rangle_{\mathbb{R}^k}) d\mu(x) \right)$$

Remark 14. 1. Key features is the factorization of $\langle n(\theta)T(x) \rangle_{\mathbb{R}^k}$.

2. Exponential forms are motivated by finding general models in which CR-efficient procedures exists.

Example 2. Binomial $p_\theta = \text{Bin}(n, \theta)$

$$\begin{aligned} p_\theta(k) &= \binom{n}{k} \theta^k (1 - \theta)^{n-k} \\ &= \binom{n}{k} \exp(k \log \theta + (n - k) \log(1 - \theta)) \\ &= \underbrace{\binom{n}{k}}_{c(k)} \exp(\underbrace{k \log \frac{\theta}{1 - \theta}}_{T(k) \cdot n(\theta)} + \underbrace{n \log(1 - \theta)}_{\Psi(\theta)}) \end{aligned}$$

Example 3. Normal $\theta = (\mu, \sigma^2) \in \mathbb{R} \times (0, \infty)$, $p_\theta = N(\mu, \sigma^2)$

$$p_{\mu, \sigma^2} = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{x^2 - 2x\mu - \mu^2}{2\sigma^2}\right)$$

Take $T(x) = \begin{pmatrix} x^2 \\ x \end{pmatrix}$, $n(\theta) = \begin{pmatrix} -\frac{1}{2\sigma^2} \\ \frac{\mu}{\sigma^2} \end{pmatrix}$ for $k = 2$, then

$$p_{\mu, \sigma^2}(x) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(\langle n(\theta), T(x) \rangle - \frac{\mu^2}{2\sigma^2}\right)$$

Example 4. *Poisson* $X = \{0, 1, \dots\}$, $p_\theta = \text{Poisson}(\theta)$, $\theta > 0$,

$$p_\theta(k) = e^{-\theta} \frac{\theta^k}{k!} = \exp(-\theta + k \log \theta) \frac{1}{k!}$$

Take $n(\theta) = \log \theta$, $T(k) = k$, $c(k) = \frac{1}{k!}$, $\Psi(\theta) = \theta$.

Definition 14. *Natural Exponential Family Define*

$$\Xi = \left\{ n \in \mathbb{R}^k \mid \int_X c(x) \exp(\langle n(\theta)T(x) \rangle_{\mathbb{R}^k}) d\mu(x) < \infty \right\}$$

A model $(P_n \mid n \in \theta)$ is called an *natural Exponential family* if $\Theta = \Xi$,

$$\frac{dP_n}{d\mu} = c(x) \exp(\langle n, T(x) \rangle_{\mathbb{R}^k} - \Psi(n))$$

for all $x \in X, n \in \Xi$.

Remark 15. A *natural Exponential family* is specified by $c(x)$, $T(x)$, and Ξ .

Lemma 4. Let $(P_n \mid n \in \Xi)$ be a 1-parameter natural Exponential form. For all $n \in \text{int}(\Xi)$

1. $\Psi'(n) = E_n[T]$
2. $\Psi''(n) = \text{var}_n[T]$

Proof. Let $n \in \text{int}(\Xi)$, and define

$$\gamma(n) = e^{\Psi(n)} = \int_X c(x) \exp(nT(x)) d\mu(x).$$

We show that γ is infinitely differentiable at n . Observe that

$$\frac{d}{dn}(c(x) \exp(nT(x))) = c(x)T(x) \exp(nT(x))$$

To use dominated convergence theorem, we want

$$\sup_t |c(x)T(x) \exp(n+t)T(x)|$$

is integrable for some $\varepsilon > 0$. But for some ε small enough, $n \neq t \in \text{int}(\Xi)$ and

$$T(x) \exp(nT(x)) \leq C \exp((n + \varepsilon)T(x))$$

for some constant $C > 0$, using that $x \leq Ce^{\varepsilon x}$ for all ε .

Thus, by using DCT

$$\begin{aligned} \frac{d}{dn} \gamma(n) &= \int_X c(x)T(x) \exp(nT(x)) d\mu(x), \\ \frac{d}{dn} \Psi(n) &= \frac{d}{dn} \log \gamma(n) = \frac{\gamma'(n)}{\gamma(n)} = E_n[T]. \end{aligned}$$

□

12 Lecture 14

Exponential Families: Assume $\Theta \subseteq \mathbb{R}$ open,

1. $p_\theta(x) = c(x) \exp\{n(\theta)T(x) - \Psi(\theta)\}$
2. Natural EF: $p_n(x) = c(x) \exp\{nT(x) - \Psi(n)\}$
3. Natural Parameter Space: $\Xi = \{t \mid \underbrace{\int_X c(x) \exp(tT(x)) d\mu(x)}_{e^{\Psi(n)}} < \infty\}$ we need to check that Ξ is an open interval.

4. $p_\theta(x)$ satisfies the regularity assumptions.

Lemma 5. $(P_n)_{n \in \Xi}$ a natural and regular EF, then for every $n \in \Xi = \Xi_0$,

1. $\varphi'(n) = E_n[T]$
2. $\varphi''(n) = \text{Var}_n(T)$

Proof. 1. Done in last Lecture

2. Recall that $\alpha(n) = \int_X c(x) \exp(nT(x)) d\mu(x) = e^{\Psi(n)}$ we had shown that α is C^∞ on this natural parameter space Ξ . as well as $\Psi'(n) = \frac{\alpha'(n)}{\alpha(n)} = E_n[T]$. Similarly,

$$\Psi''(n) = \frac{\alpha''(n)}{\alpha(n)} - \frac{\alpha'(n)^2}{\alpha(n)^2} = \int_X c(x) T(x)^2 \exp(nT(x)) d\mu(x) - E_n[T]^2 = E_n[T^2] - E_n[T]^2 = \text{Var}_n(T)$$

□

Example 5. Often $T(x) = x$, like with the Poisson, Normal, etc.

1. $P_\theta = \text{Bin}(n, \theta)$, $T(k) = k$, $n = \log \frac{\theta}{1-\theta}$

$$\text{Recall } p_\theta(k) = \binom{n}{k} \theta^k (1-\theta)^{n-k} = \binom{n}{k} \exp\left(k \log \frac{\theta}{1-\theta} + n \log(1-\theta)\right)$$

$$\begin{aligned} \Psi(n) &= -n \log(1-\theta) = -n \log\left(1 - \frac{e^n}{1+e^n}\right) \\ \theta &= \frac{e^n}{1+e^n} \\ &= -n \log\left(\frac{1}{1+e^n}\right) = n \log(1+e^n) \end{aligned}$$

Hence,

$$\Psi'(n) = n \frac{e^n}{1+e^n} = n\theta = E_n[T] = \text{Mean of Bin}(n, \theta)$$

$$\Psi(n) = \text{Var}_n[T] = n\theta(1-\theta)$$

2. $P_\lambda = \text{Poisson}(\lambda)$,

$$P_\lambda(k) = e^{-\lambda} \frac{\lambda^k}{k!} = \frac{1}{k!} \exp\left(\underbrace{k}_{T(k)} \underbrace{\log \lambda}_n - \underbrace{\lambda}_{\Psi(n)}\right)$$

with $\Psi(n) = e^n$, then $\Psi' = \Psi'' = e^n$. Hence,

$$E_n[T] = e^n = \lambda, \text{Var}_n(T) = e^n = \lambda$$

Theorem 9. MLEs in Internal EF and UMVUE Estimators Let $(P_\theta)_{n \in \Xi}$ be a natural 1-parameter regular EF, then

1. If a unique MLE \hat{n}_{MLE} exists then $(\Psi')^{-1}(T) = \hat{n}_{MLE}$
2. Define $\rho(n) = E_n[T]$ where T is the UMVUE for $\rho(n)$

Proof. 1.

$$\begin{aligned} \hat{n}_{MLE} &= \hat{n} = \arg \max_{n \in \Xi} p_n(x) \\ &= \arg \max_n \log p_n(x) \\ &= \arg \max_n nT(x) - \Psi(n) \end{aligned}$$

Therefore,

$$\left. \frac{d}{d\mu} (nT(x) - \Psi(n)) \right|_{n=\hat{n}} = 0 \implies T(x) = \Psi'(\hat{n})$$

Moreover, $\Psi''(n) = \text{Var}_n(T) > 0$ results from inverting Ψ' .

2. Recall CR lower bound $\text{Var}_n(S) \geq \rho'(n)^2 I'(n)$ for any unbiased estimator S . It holds that

$$\rho'(n) = \frac{d}{dn}(E_n[T]) = \Psi''[n]$$

where $I(n) = E_n[\log p_n(x)^2] = -E_n[\log p_n(x)']$ Then

$$(\log p_n(x))' = (nT(x) - \Psi(n))' = T(x) - \Psi'(n) = T(x) - E_n[T(x)]$$

$$E_n[(\log p_n(x))'^2] = \text{Var}_n(T) = \Psi''(n).$$

Therefore the CR-bound $\text{Var}_n(S) \geq \Psi(n) = \text{Var}_n(T)$

□

13 Generalized Linear Models (GLMs)

Linear models + Exponential Families → GLMs.
 establish relations, $x_i \rightarrow y_i, X\beta + \varepsilon = Y$ flexible classes continuous (normal) discrete (Poisson) Binomial (Bernoulli)
 Note, General (X, F) are allowed, but Θ will still be an open set in \mathbb{R}^p

Example 6. Suppose we have binary data, $Y_i \in \{0, 1\}$, $i = 1, \dots, n$

1. Covariates $X_i \in \mathbb{R}^p$
2. Logistic Regression for $\beta \in \mathbb{R}^p$,

$$p_\beta(Y_i = 1|X_i) = \frac{\exp(X_i^T \beta)}{1 + \exp(X_i^T \beta)}$$

Equivalently,

$$X_i^T \beta = \log \left(\frac{P_\beta(Y_i = 1|X_i)}{1 - P_\beta(Y_i = 1|X_i)} \right) = \log \text{it}(P_\beta(Y_i = 1|X)).$$

Example 7. Poisson Regression

1. $Y_i \sim \text{Poisson}(\lambda_i)$, $\lambda_i > 0$.
2. $\log \lambda_i = X_i^T \beta$
3. $X_i \in \mathbb{R}^p$, $\beta \in \mathbb{R}^p$
4. Used for count data

Definition 15. Generalized Linear Models We have the data $(X_i, Y_i) \in \mathbb{R}^p \times \mathbb{R}$ In a GLM,

$$dP_{n_i}^{Y_i}(y_i) = c(y_i) \exp(n_i y_i - \Psi(n_i))$$

where n_i and x_i are linked through some link function

$$g: \mathbb{R} \rightarrow \mathbb{R},$$

$$g(E_{n_i}[Y_i]) = X_i^T \beta.$$

Intuition, $n_i \iff E_{n_i}[Y_i] = \mu_i \iff X_i^T \beta$ with $g' > 0$.

When $g(E_{n_i}[Y_i]) = n_i$, then g is called the canonical link, or the natural link function.

Remark 16. 1. Under the canonical link,

$$p_\beta^{Y_i}(Y_i) = c(y_i) \exp(Y_i X_i^T \beta - \Psi(X_i^T \beta)).$$

2. Link function links the linear predictors X_i to the mean of the outcome.

$$E[Y_i|X_i, \beta] = g^{-1}(X_i^T \beta)$$

Lecture 15: GLMs & Model Selection

Natural Exponential Family (Nat. EF)

$$P_\eta(x) = c(x) \exp(\eta T(x) - \Psi(\eta)), \quad \eta \in \mathbb{R}$$

- What is the Fisher information?
- Is the Cramer-Rao lower bound attained?

$$\mathbb{E}_\eta[T] = \Psi'(\eta)$$

- What happens if $T = \text{const. a.s.}$?

Generalized Linear Models (GLM)

- Parameter: $\beta \in \mathbb{R}^p$
- Design matrix: $X \in \mathbb{R}^{n \times p}$
- Data: $Y_i \sim c(y) \exp(y X_i^\top \beta - \Psi(X_i^\top \beta))$
-

$$g(\mathbb{E}[Y_i | X_i, \beta]) = X_i^\top \beta$$

- (Natural) link function: $g : \mathbb{R} \rightarrow \mathbb{R}$

Example: Logistic Regression

$$Y_i \in \{0, 1\}, \quad \mathbb{E}[Y_i] = P(Y_i = 1) = p_i$$

- Natural link function:

$$g(p_i) = \log \frac{p_i}{1 - p_i}, \quad g^{-1}(X_i^\top \beta) = \frac{\exp(X_i^\top \beta)}{1 + \exp(X_i^\top \beta)}$$

•

$$\begin{aligned} P_\beta(Y_i) &= p_i^{Y_i} (1 - p_i)^{1 - Y_i} \\ &= \exp(Y_i \log p_i + (1 - Y_i) \log(1 - p_i)) \\ &= \exp\left(Y_i \log \frac{p_i}{1 - p_i} + \log(1 - p_i)\right) \\ &= \exp(Y_i X_i^\top \beta - \Psi(X_i^\top \beta)), \quad \Psi(X_i^\top \beta) = \log(1 + \exp(X_i^\top \beta)) \end{aligned}$$

MLE in Logistic Regression

Goal: Find $\hat{\beta} \in \arg \max_{\beta \in \mathbb{R}^p} P_\beta(Y)$.

$$\begin{aligned} \log p_\beta(Y) &= \sum_{i=1}^n \log p_i^{Y_i} (1 - p_i)^{1 - Y_i} \\ &= \sum_{i=1}^n Y_i \log \frac{p_i}{1 - p_i} + \log(1 - p_i) \\ &= \sum_{i=1}^n y_i \underbrace{\eta_i}_{X_i^\top \beta} + \log \frac{1}{1 + e^{\eta_i}} \\ \nabla_\beta \log p_\beta(Y) &= \sum_{i=1}^n Y_i \nabla_\beta \eta_i - \frac{1}{1 + e^{\eta_i}} e^{\eta_i} \nabla_\beta \eta_i \\ &= \sum_{i=1}^n \nabla_\beta \eta_i \left(Y_i - \underbrace{\frac{e^{\eta_i}}{1 + e^{\eta_i}}}_{p_i} \right) = \sum_{i=1}^n (Y_i - p_i) X_i^\top \\ &= X^\top (y - p) \in \mathbb{R}^p \end{aligned}$$

MLE in Logistic regression,

$$X^T(Y - p) \neq 0$$

where $p = (p_1, \dots, p_n)$ is $p_i = \frac{e^{X_i^T \beta}}{1 + e^{X_i^T \beta}}$

Caution, $\hat{\beta}$ is not always defined.

We know $E[Y_i] = P(Y_i = 1) = \frac{e^{X_i^T \beta}}{1 + e^{X_i^T \beta}}$.

Remark 17. $\hat{\beta}$ the MLE in GLMs are not given in closed form, and have to be computed using optimization methods (e.g. gradient decent, Newton's Method)

Lemma 6. Score and FI in GLMs Consider a GLM with canonical links. Then,

$$\begin{aligned}\nabla_{\beta} \log p_{\beta}(Y) &= \sum_{i=1}^n Y_i - \Psi'(X_i^T \beta) X_i^T \\ I(\beta) &= \sum_{i=1}^n \Psi''(X_i^T \beta) X_i X_i^T \in \mathbb{R}^p\end{aligned}$$

Proof.

$$\begin{aligned}\log p_{\beta}(Y) &= \sum_{i=1}^n \log c(y_i) + y_i X_i^T \beta - \Psi(X_i^T \beta) \\ \implies \nabla_{\beta} \log p_{\beta}(Y) &= \sum_{i=1}^n Y_i X_i^T - \Psi'(X_i^T \beta) X_i^T = \sum_{i=1}^n (Y_i - \psi'(X_i^T \beta)) X_i^T\end{aligned}$$

Once checks that

$$\frac{\partial^2}{\partial \beta_k \partial \beta_{\ell}} \log p_{\beta}(Y) = \sum_{i=1}^n -\Psi'' \Psi''(X_i^T \beta) (x_i)_k (x_i)_{\ell}$$

Since $I(\beta) = E_{\beta}[-\nabla_{\beta}(\log p_{\beta}(x))]$, the claim follows. \square

Remark 18. 1. In natural ER and in GLMs, the curvature of $\log p_{\beta}$ is independent of the data!

2. If $\hat{\beta}$ MLE exists in \mathbb{R}^p and $I(\beta)$ is positive definite, then $\hat{\beta}$ is unique (not clear in general)

13.1 Model Selection

Setting: Suppose we observe $Y \in \mathbb{R}^n$ of the form, $Y = \mu + \varepsilon$ with unknown $\mu \in \mathbb{R}^n, \varepsilon \sim N(0, \sigma^2 I_n)$, σ^2 unknown.

For $k = 1, \dots, K \leq n$, suppose we have linear models for μ ,

$$X^{(k)} \beta^{(k)}, \beta^{(k)} \in \mathbb{R}^k$$

where $X^{(k)} \in \mathbb{R}^{n \times k}$ with rank k .

Example 8. Full design matrix,

$X^{(k)}$ = first k columns of X where X is a k by n matrix

We can ask ourselves, what's the best model.

For example,

$$\hat{\beta}^{(k)} = ((X^{(k)})^T X^{(k)})^{-1} X^{(k)T} Y = \arg \min_{\beta^{(k)} \in \mathbb{R}^k} \|Y - X^{(k)} \beta^{(k)}\|$$

Which has the least MSE.

$$\|\underbrace{\mu - X^{(k)} \hat{\beta}^{(k)}}_{\hat{\mu}^{(k)}}\|^2$$

We calculate,

$$E\|\mu - \hat{\mu}^{(k)}\|_{\mathbb{R}^k}^2 = (\mu - E[\hat{\mu}^{(k)}])^2 + E[(\mu^{(k)} - E[\hat{\mu}^{(k)}])^2] - \underbrace{2 E[\langle \mu - E[\hat{\mu}^{(k)}], \hat{\mu}^{(k)} - E[\hat{\mu}^{(k)}] \rangle]}_{=0}$$

Moreover,

$$E[\hat{m}u^{(k)}] = E[\Pi^{(k)}Y] = E[\Pi^{(k)}(\mu + \varepsilon)] = \Pi^{(k)}\mu$$

So then,

$$\begin{aligned} \hat{\mu}^{(k)} - E[\hat{\mu}^{(k)}] &= \Pi^{(k)}\varepsilon \\ \implies E\|\mu - \hat{\mu}^{(k)}\|_{\mathbb{R}^k}^2 &= \|(I_n - \Pi^{(k)})\mu\|^2 + \underbrace{E[\Pi^{(k)}\varepsilon]^2}_{\sigma^2\chi^2(k)} \end{aligned}$$

and all of this is equal to

$$= \|(I_n - \Pi^{(k)})\mu\|^2 + k\sigma^2 = \text{BIAS} + \text{VARIANCE with } k$$

How well is μ approximated by $\text{col}(X^{(k)})$ We'd like to pick

$$\hat{k} = \arg \min_{k=1, \dots, K} \underbrace{\|(I_n - \Pi^{(k)})\mu\|}_{\text{Unknown}} + k\sigma^2$$

To estimate the first term, we consider RSS,

$$\begin{aligned} E\|\underbrace{Y - X^{(k)}\hat{\beta}^{(k)}}_{\text{Data Driven}}\|^2 &= E\|(I_n - \Pi^{(k)})(\mu + \varepsilon)\|^2 \\ &= E\|(I_n - \Pi^{(k)})\mu\|^2 + E\|(I_n - \Pi^{(k)})\varepsilon\|^2 + E[\underbrace{\langle (I_n - \Pi^{(k)})\mu, (I_n - \Pi^{(k)})\varepsilon \rangle}_{\text{detm}} \underbrace{\quad}_{E[\cdot]=0}] \\ &= \|(I_n - \Pi^{(k)})\mu\|^2 + \sigma^2(n - k) \end{aligned}$$

Which implies

$$\|Y - X^{(k)}\hat{\beta}^{(k)}\|^2 - \sigma^2(n - k) + \sigma^2k$$

is unbiased risk estimator for risk, $E\|\mu - \hat{\mu}^{(k)}\|^2$

Method (Mallow's Cp)

Pick $\hat{k} = \arg \min_{k=1, \dots, K} \|Y - X^{(k)}\hat{\beta}^{(k)}\|^2 + 2\sigma^2k$

Next time we will generalize this idea to Akaike's Information Criterion. (AIC).